internet protocol suite

javed shaikh

/agenda

- Introduction
  - Networking technology: Ethernet
  - IP addressing
  - Address Resolution Protocols
- Transport layer
  - Transmission Control Protocol (TCP)
  - User Datagram Protocol (UDP)
- Network layer
  - Internet Protocol (IPv4)
  - Internet Control Message Protocol (ICMP)

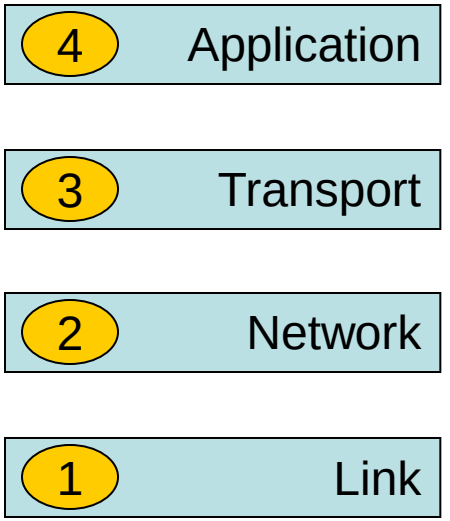

/references

- TCP/IP – Behrouz Forouzan, *McGraw Hill*
- TCP/IP Guide – Charles M. Kozierok, *No Starch Press* www.tcpipguide.com
- TCP/IP Illustrated Volume 1 (The Protocols) – W. Richard Stevens, *Addison-Wesley*

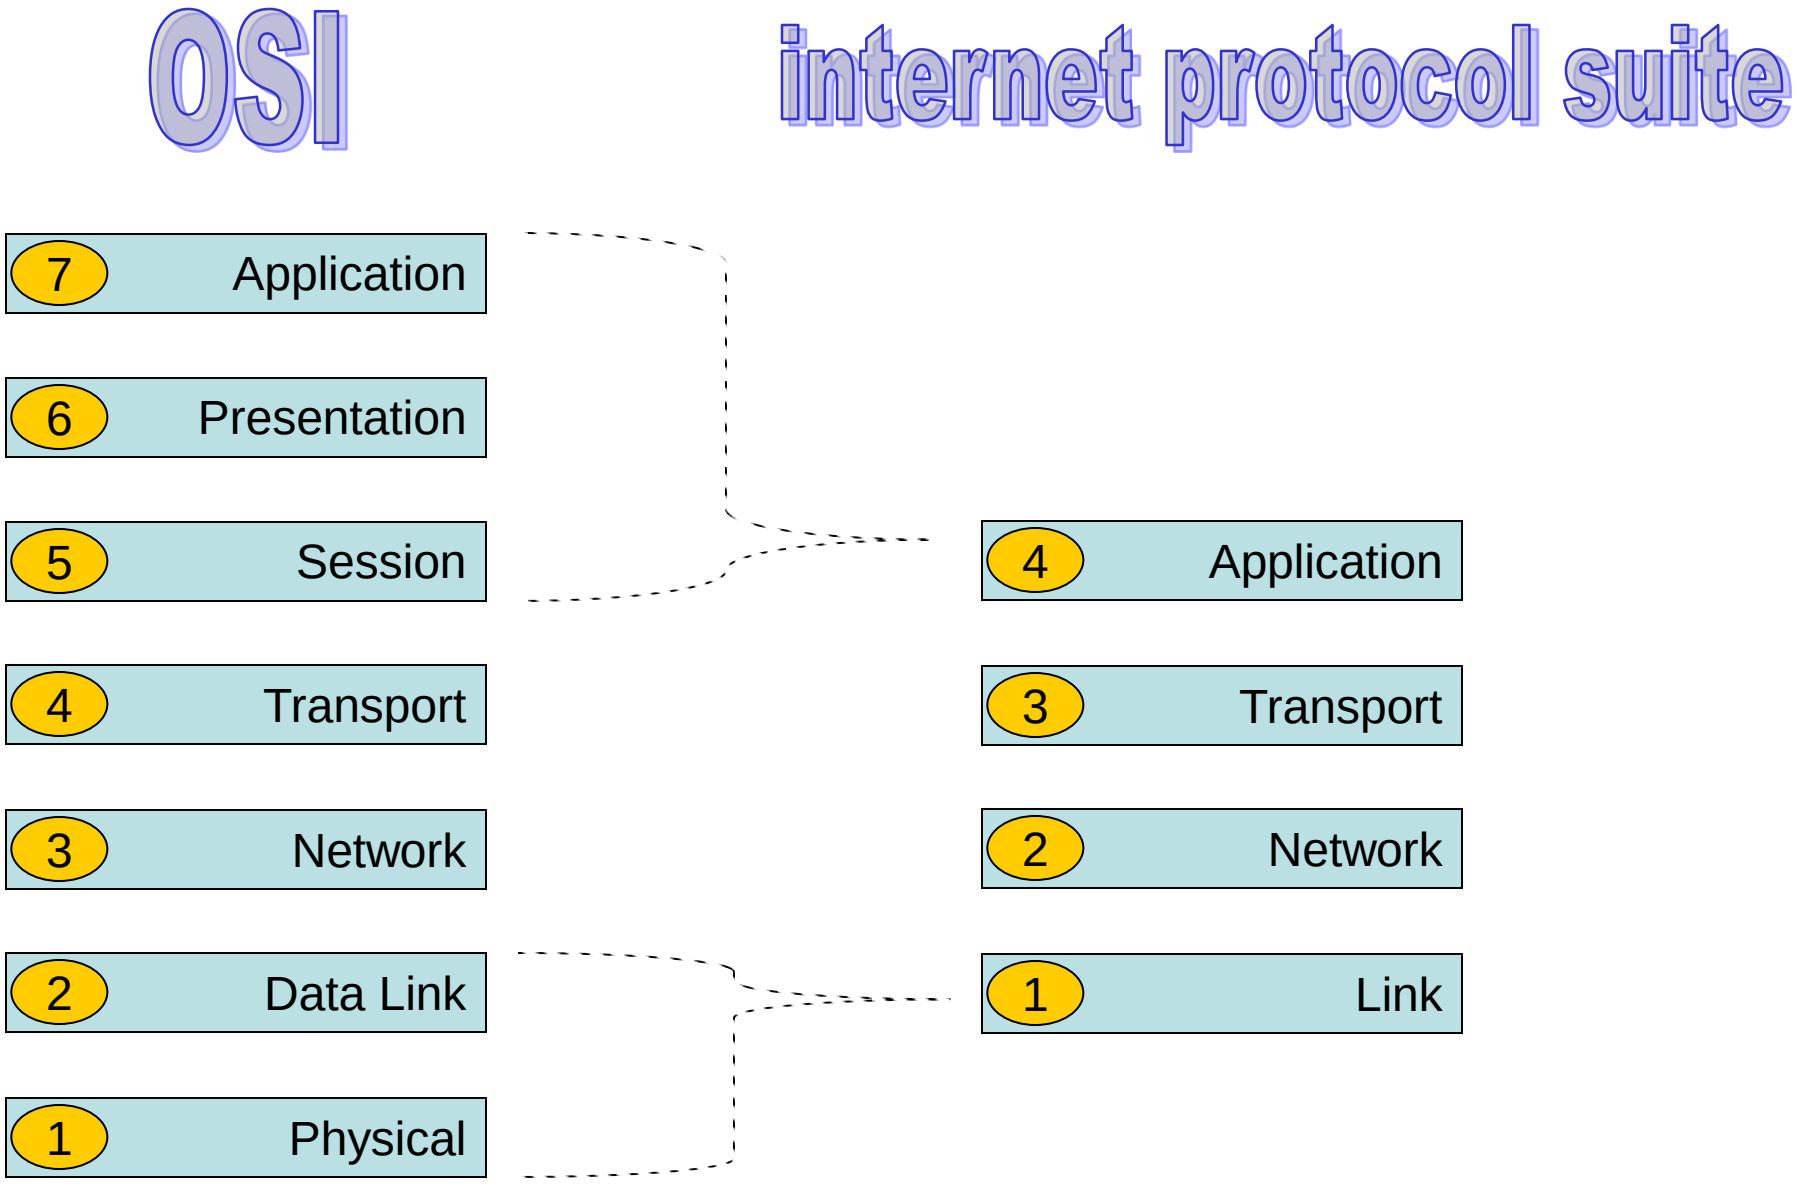

/overview/stack

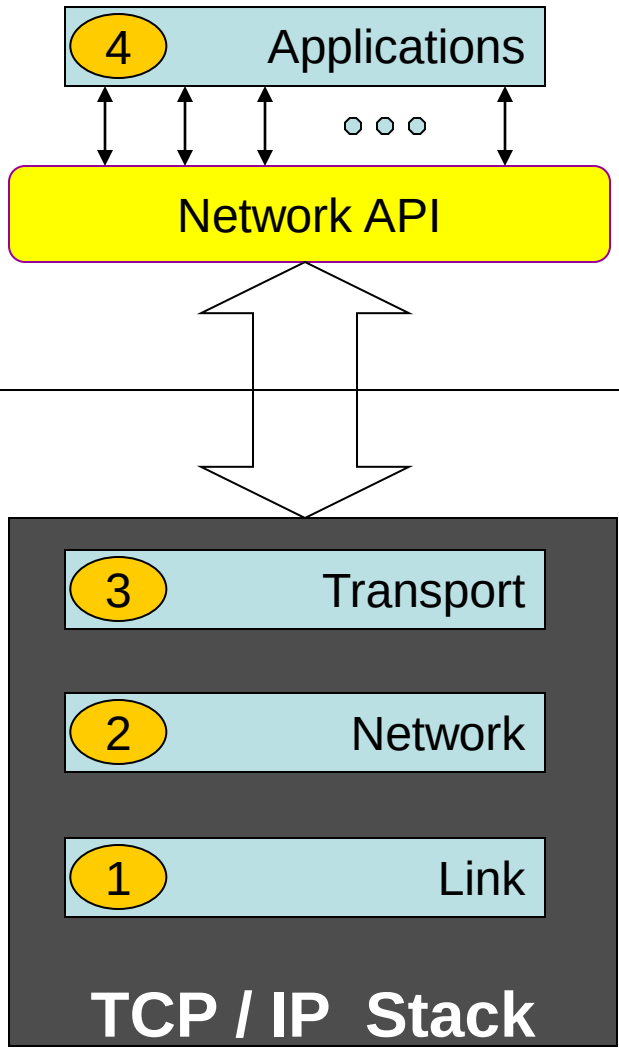

telnet, ftp, http, e-mail, chat, etc.

AT&T Transport Layer Interface, Berkeley sockets

user

kernel

TCP, UDP → protocol drivers

IP, ICMP → protocol drivers

Network cable, Network interface card, device driver

/overview/applications

User Applications	Administrative Utilities
telnet	hostname
ftp	ping
talk	traceroute
rlogin	arp
rsh	ifconfig
mozilla, iexplorer (http)	netstat
mail	

Link Layer

- Networking technologies
  - Token ring
  - ATM
  - FDDI
  - **Ethernet**
  - RS-232 serial line

/link layer/ethernet

- Inventors – DEC, Intel, Xerox
- Access method – CSMA/CD
- Speed – 10 (E) / 100 (FE) /
  1000 (GigE) / 10000 (10G)
- Address – 48 bits
- RFC 894 (Ethernet) *Vs* RFC 1042 (IEEE 802)
- Requirements from Internet host:
  - Must tx/rx Ethernet packets (default)
  - Should be able to rx IEEE 802 packets intermixed with Ethernet
  - May be able to send IEEE 802 packets. If sending both types, the type of packet sent must be configurable

/link layer/ethernet/differences

IEEE 802

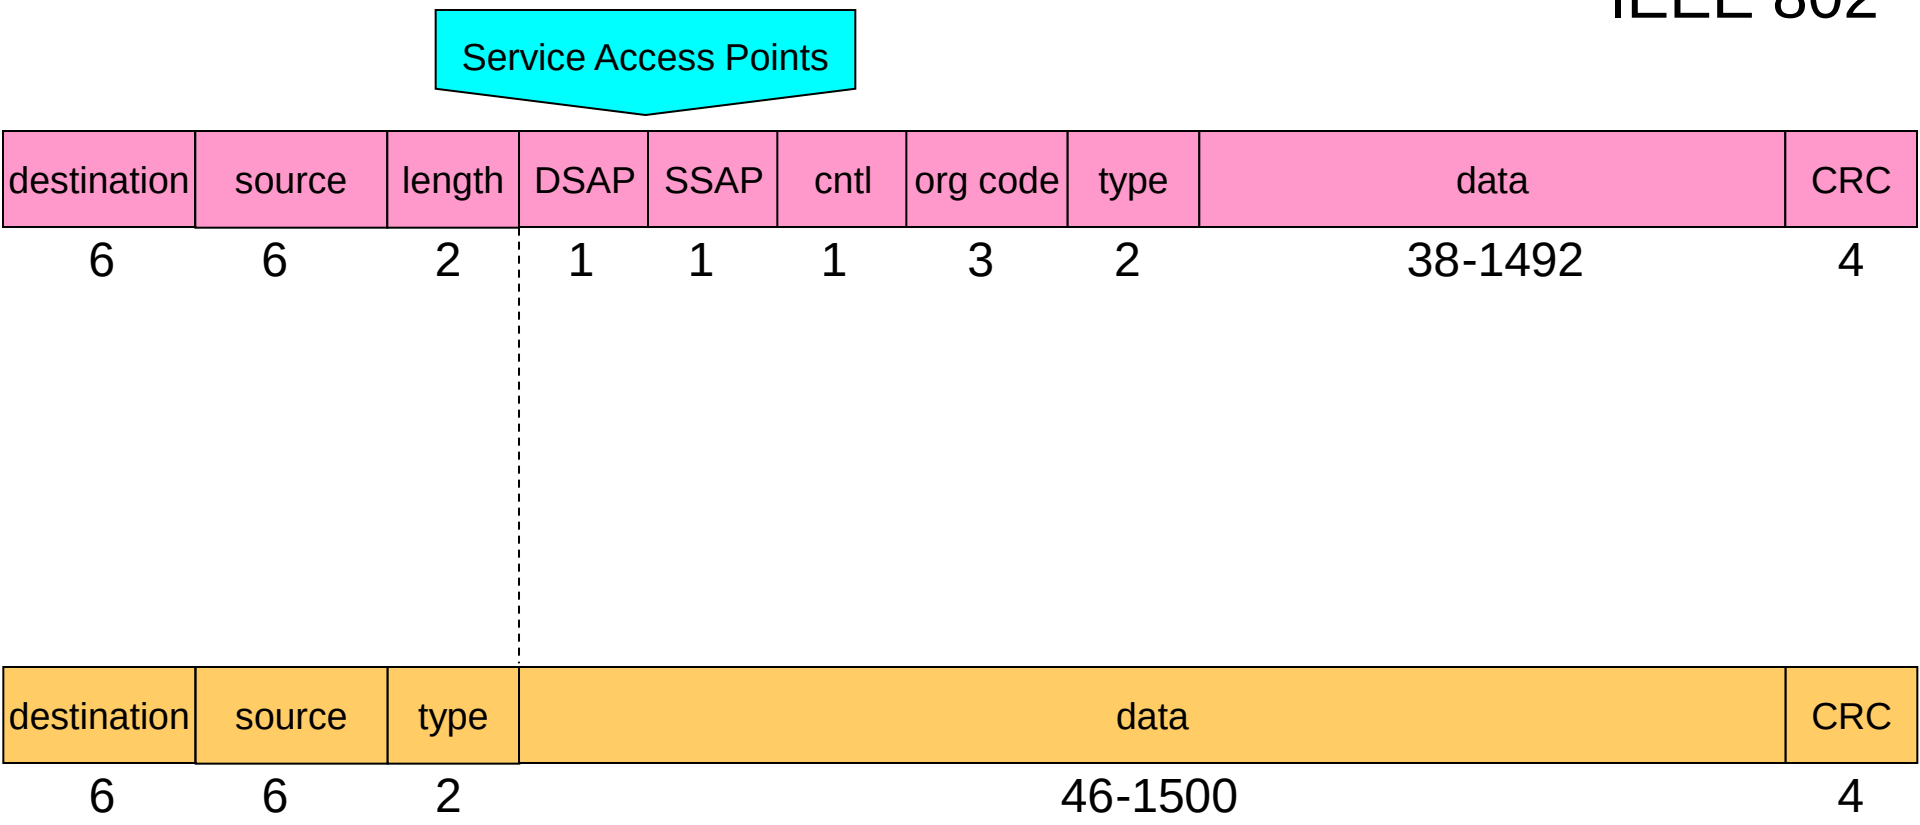

Ethernet

Minimum 64 byte frame required…

- Every interface must have a unique IP address
- 32 bit, dotted-decimal notation
- 3 types
  - *Unicast, broadcast, multicast*
- 5 classes

A	0.0.0.0	127.255.255.255
B	128.0.0.0	191.255.255.255
C	192.0.0.0	223.255.255.255
D	224.0.0.0	239.255.255.255
E	240.0.0.0	247.255.255.255

/network layer/ip/addressing

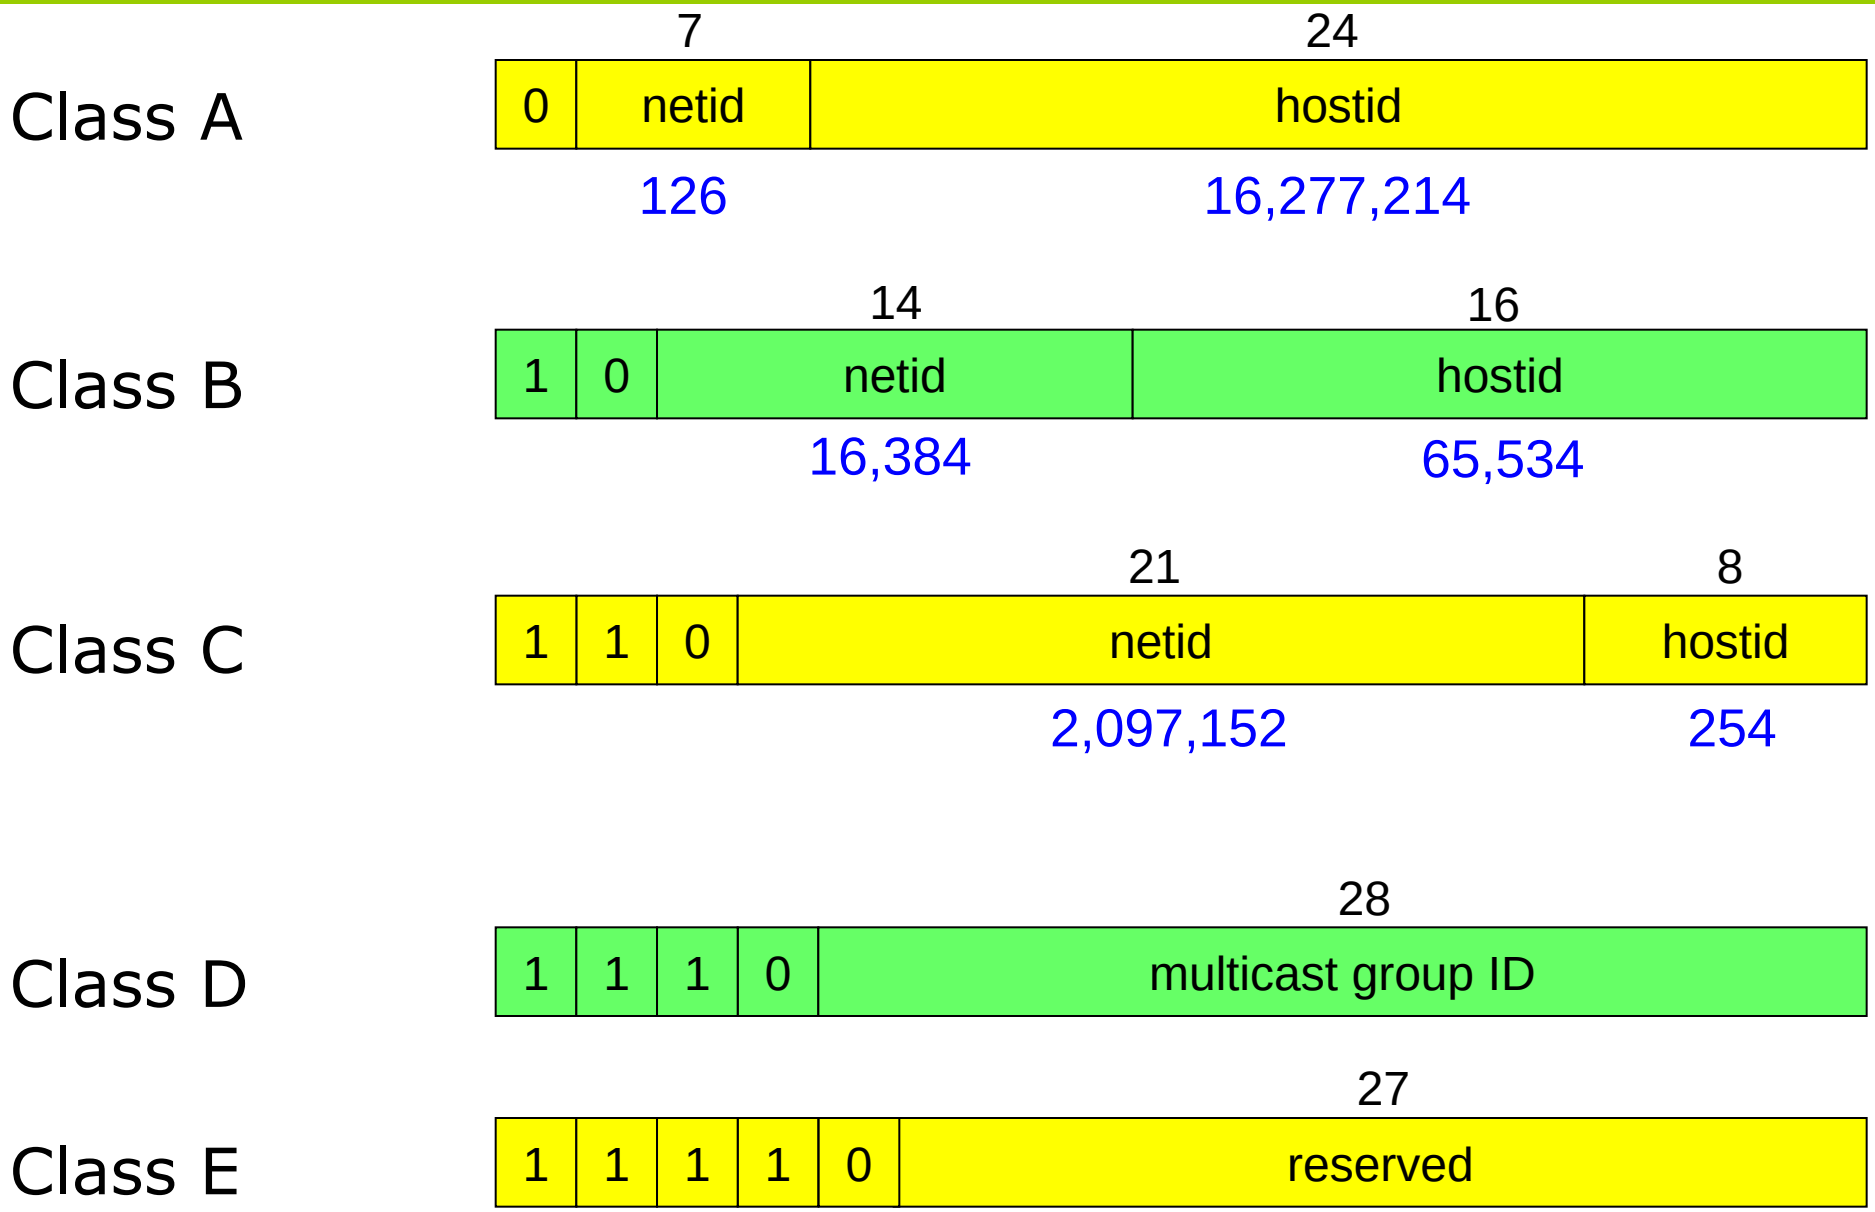

/network layer/ip/addressing

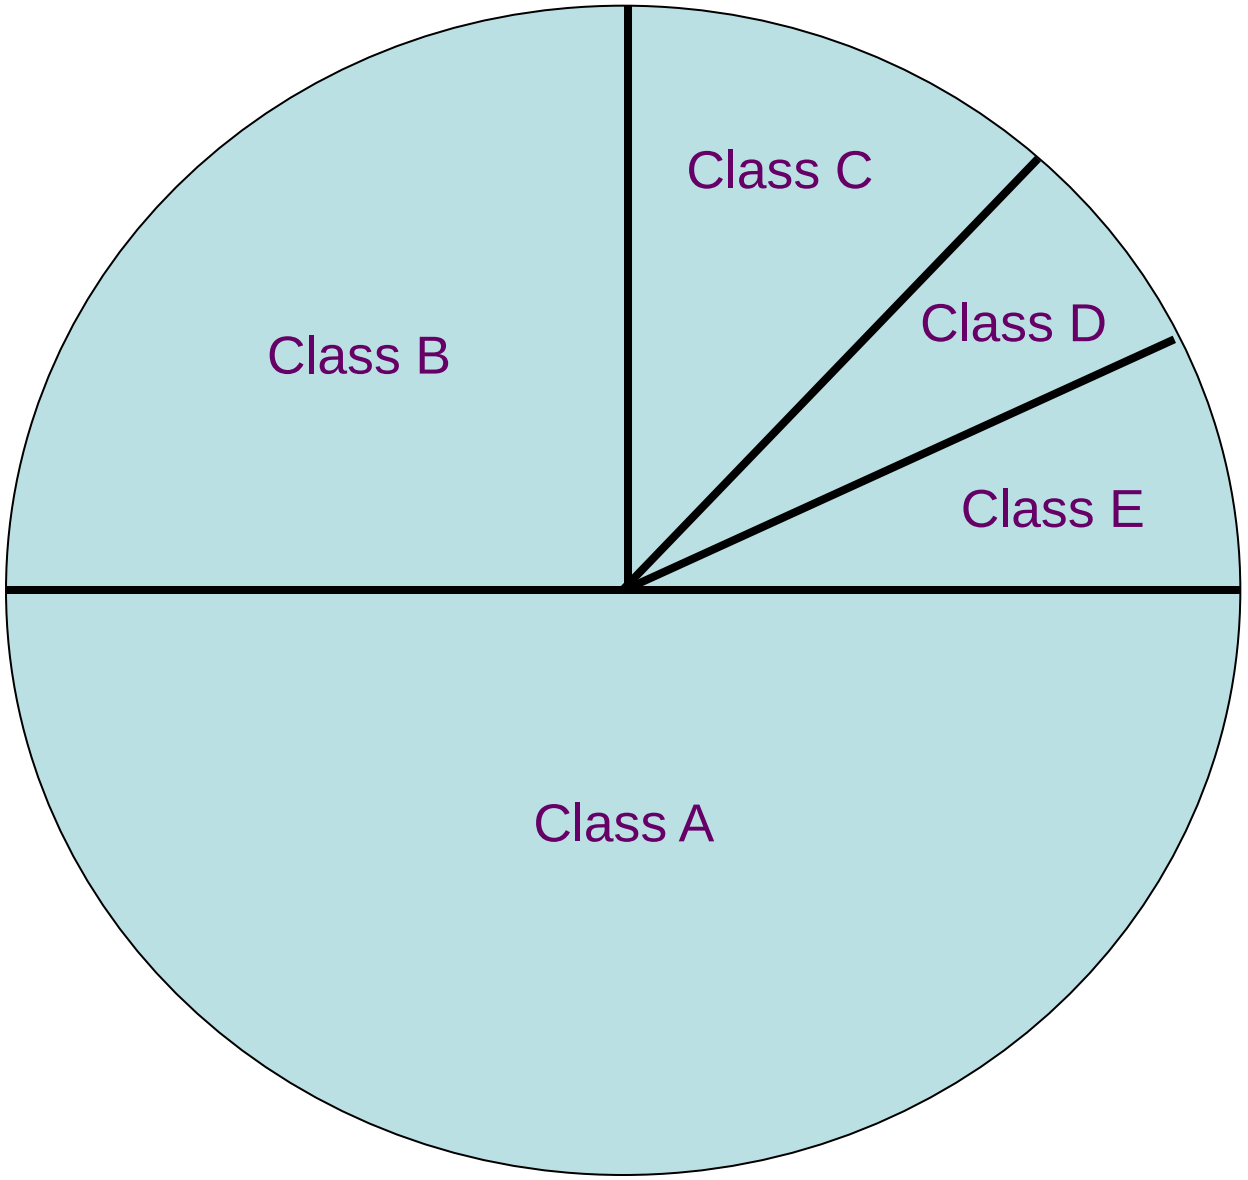

/network layer/ip/addressing/reserved

Start	End	Comment
0.0.0.0	0.255.255.255	-
10.0.0.0	10.255.255.255	Class A private addr block
127.0.0.0	127.255.255.255	Loopback addr block
128.0.0.0	128.0.255.255	-
169.254.0.0	169.254.255.255	Class B private addr block for auto addr allocation
172.16.0.0	172.31.255.255	Private addr block
191.255.0.0	191.255.255.255	-
192.0.0.0	192.0.0.255	-
192.168.0.0	192.168.255.255	Class C private addr block
223.255.255.0	223.255.255.255	-

- Class C IP address → 192.168.71
  (254 hosts possible)
- Network Mask → 255.255.255.0
  Network Mask specifies the netid
- Valid hostid → 1 to 254
- Broadcast Address → 192.168.71.255

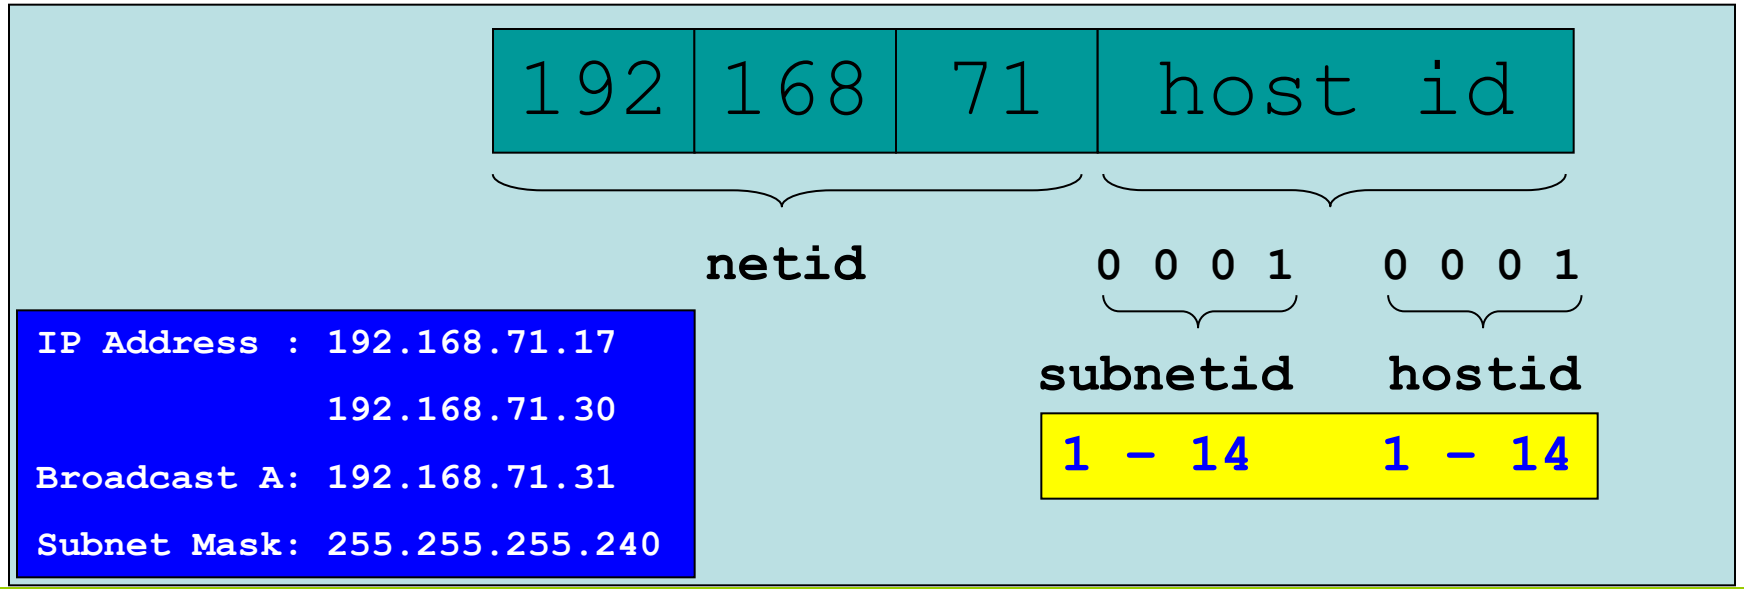

/network layer/ip/addressing/special addrs

Netid	Subnetid	Hostid	Valid as
0		0	src
0		>0	src
127		**any**	**src / dest**
-1		-1	dest
netid		-1	dest
netid	subnetid	-1	dest
netid	-1	-1	dest

/network → link layer/address mapping/introduction

- A link can be used by any network layer protocol
- An Ethernet link has its own addressing scheme and uses the same to send a frame to destination host
- IP addresses make sense only to network layer and above
- Hence a mechanism to map IP address to a MAC address is required
  - ARP : Address Resolution Protocol
- Similarly a mechanism to obtain an IP address for a MAC address is
  - RARP : Reverse Address Resolution Protocol

/network → link layer/arp

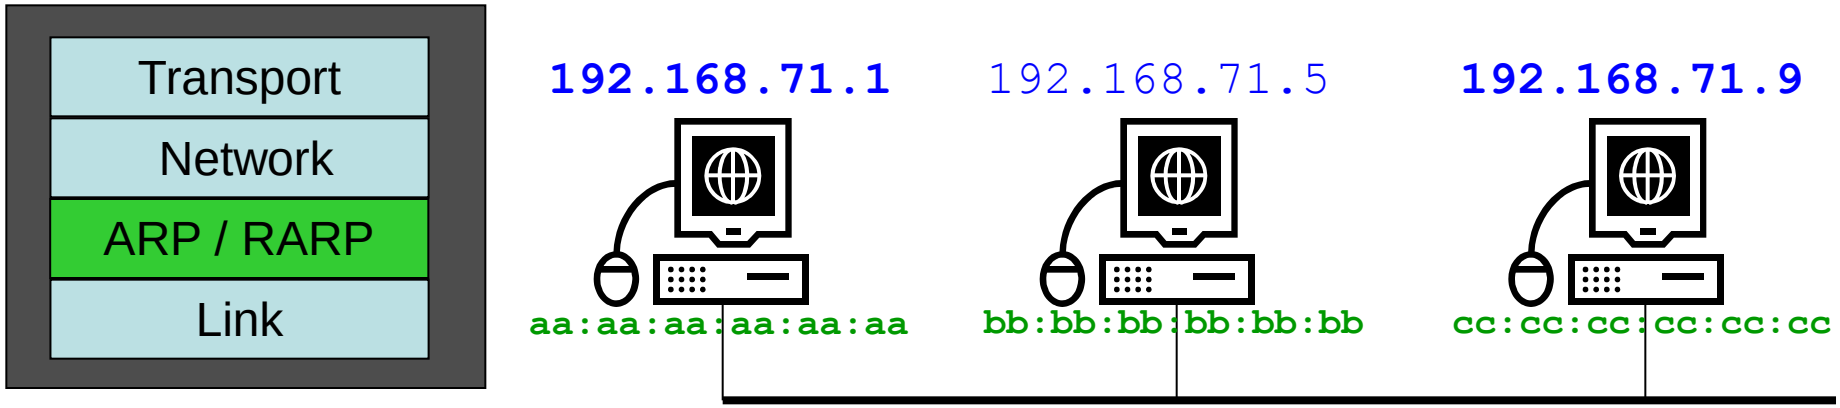

source	destination	type	data	CRC
6	6	2	46-1500	4

R/ARP Packet

hard type	proto type	HAL	PAL	op	src ether addr	src IP addr	dest ether addr	dest IP addr
2	2	1	1	2	6	4	6	4

aa:aa:aa:	**ff:ff:ff:…**	0806	1	800	6	4	1	aa:aa:aa…	192.168. 71.1	ff:ff:ff:…	192.168. 71.9	CRC

cc:cc:cc:	**aa:aa:aa:**	0806	1	800	6	4	2	cc:cc:cc:..	192.168. 71.9	aa:aa:a…	192.168. 71.1	CRC

/network → link layer/arp/details

- ARP request is broadcast and reply is unicast
- ARP cache
  - Timeout = 1 to 20 mins. For Linux check:

    /proc/sys/net/ipv4/neigh/eth0/gc_stale_time
- ARP request to non-existent host
- Proxy ARP (*For hosts on different networks at the data link layer level, but on the same IP network or subnet*)
- Gratuitous ARP
  - Detects duplication of IP address
  - Updating of older hardware address
- ARP is part of the kernel's TCP/IP implementation

/network → link layer/arp/assignments

- Use *arp* command on linux system with following options:
  - a : to display all the entries in ARP cache
  - d : to delete an entry from ARP cache
  - s : to add an entry in ARP cache
- Trap and display ARP packets using *tcpdump* (requires root privilege)
  - For a host which does not have its entry in the arp cache
  - For a host which has an entry in the arp cache

/network → link layer/rarp

- A computer acquires its IP address from a file stored on it during bootstrapping procedure. What about a computer that is diskless?

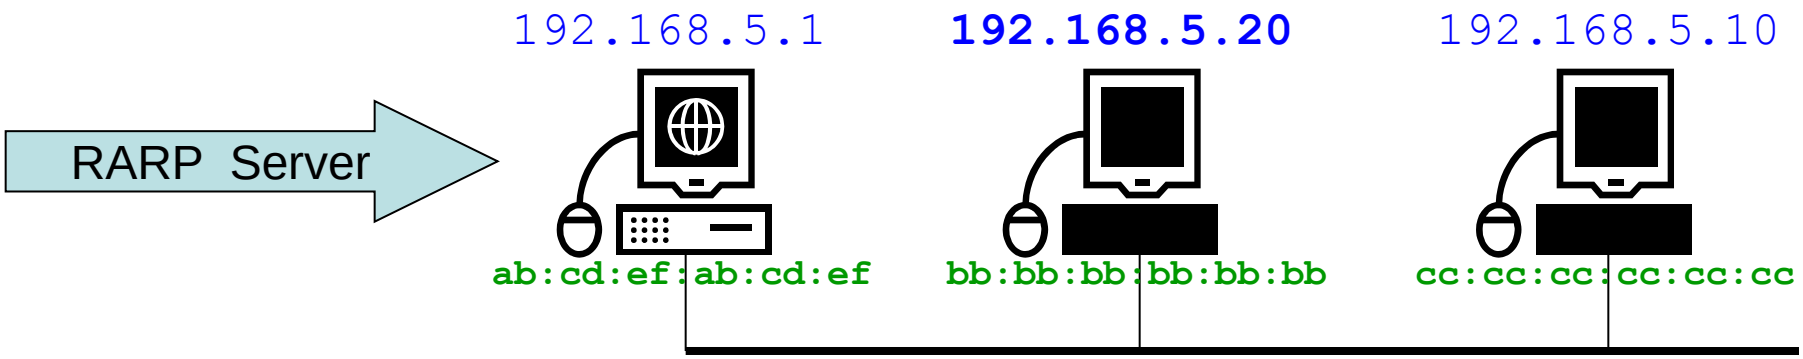

bb:bb:bb:	**ff:ff:ff:…**	8035	1	800	6	4	3	bb:bb:bb…	-	ff:ff:ff…	-	CRC

ab:cd:ef:	**bb:bb:bb..**	8035	1	800	6	4	4	ab:cd:ef…	192.168.5.1	bb:bb:bb…	192.168.5.20	CRC

- RARP request is broadcast and reply is unicast
- RARP requests are sent as hardware-level broadcasts, hence they are not forwarded by routers
- Multiple RARP servers need to be maintained to provide redundancy
- RARP server is implemented as user process. Also this implementation is tied to the system
- Made obsolete by Bootstap Protocol (BOOTP) which is now replaced by Dynamic Host Configuration Protocol (DHCP)

Terminology

/terms

- Protocol – A *networking protocol* defines a set of rules and messages that enable software and hardware in networked devices to communicate effectively

- Mbps, MBps

- Frame, Segment, Datagram, Packet, Data

Transport Layer

/transport layer

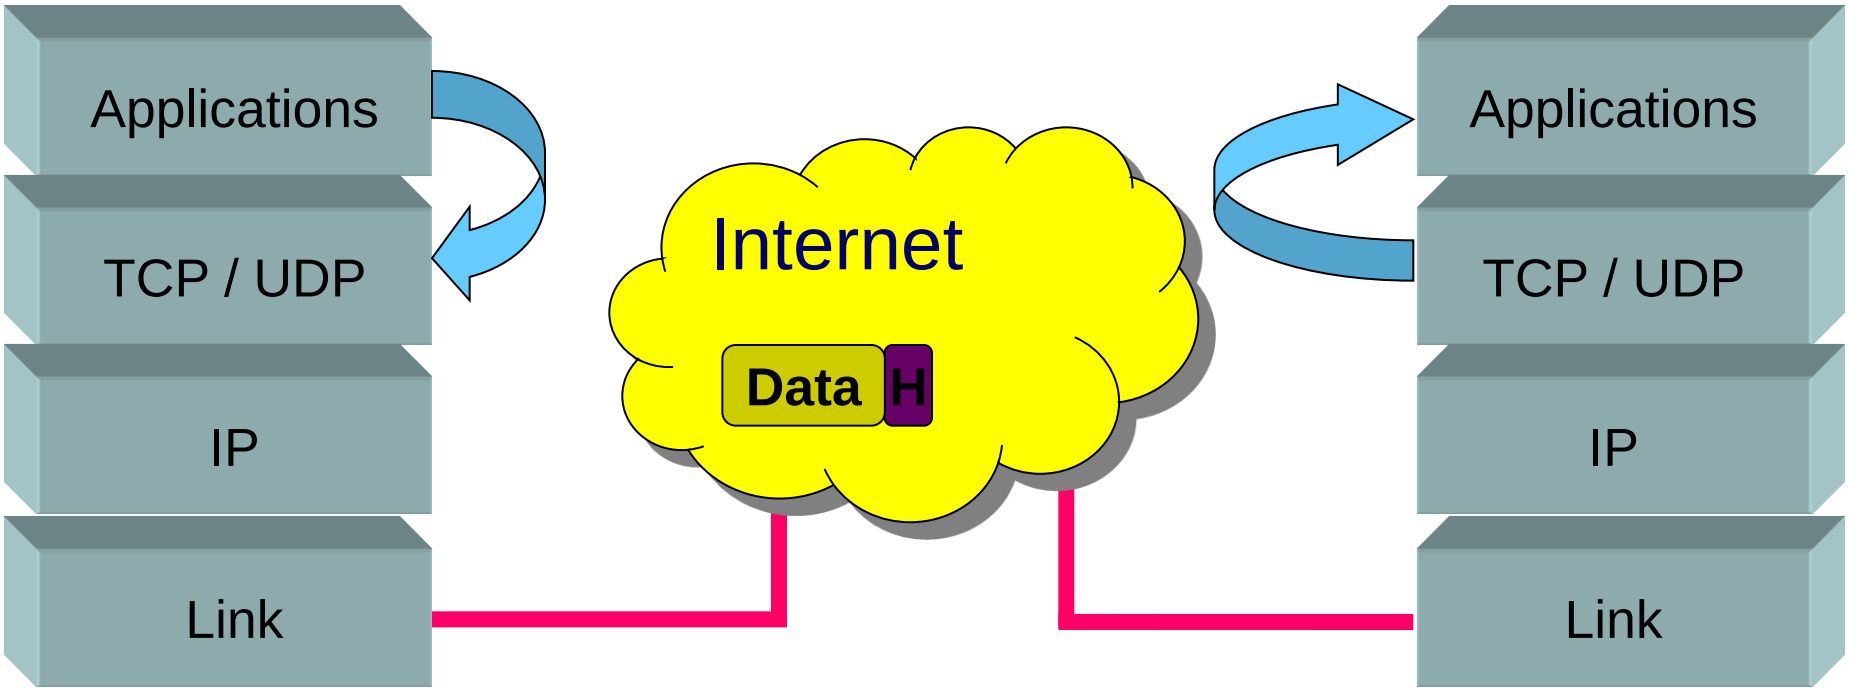

/transport layer

- Provides a flow of data between two hosts for the application layer above
- ***TCP:***
  - Provides reliable flow of data between two hosts
  - Divides data passed to it from application into appropriate sized chunks for the network layer below
  - Acknowledges received packets
  - Sets timeouts for acks, etc.
- ***UDP:***
  - Sends packets of data called **datagrams** from one host to another, but gives no guarantee that it will reach
  - Reliability if desired must be added by the application layer

/transport layer/encapsulation

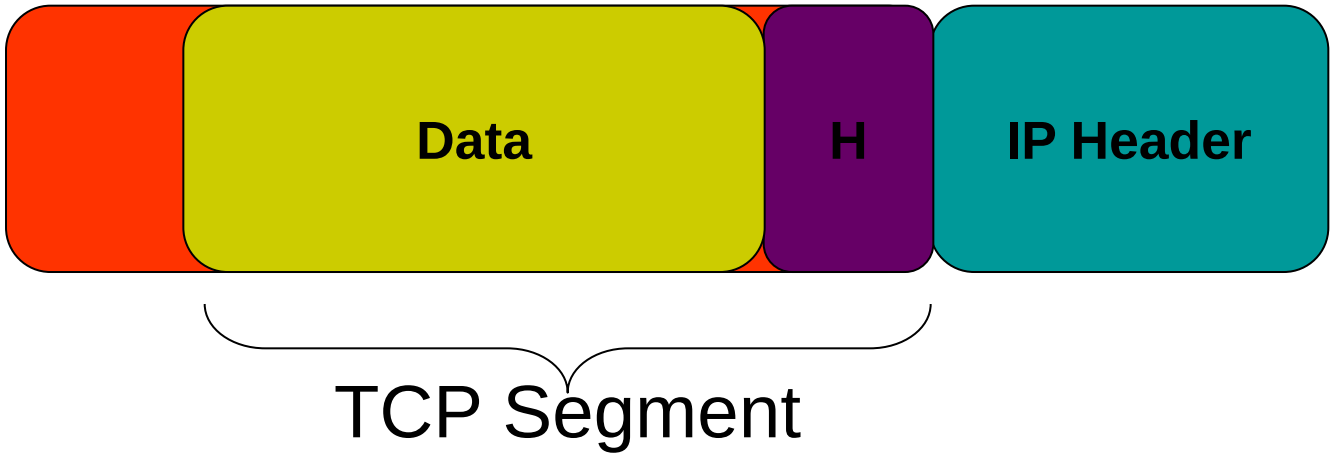

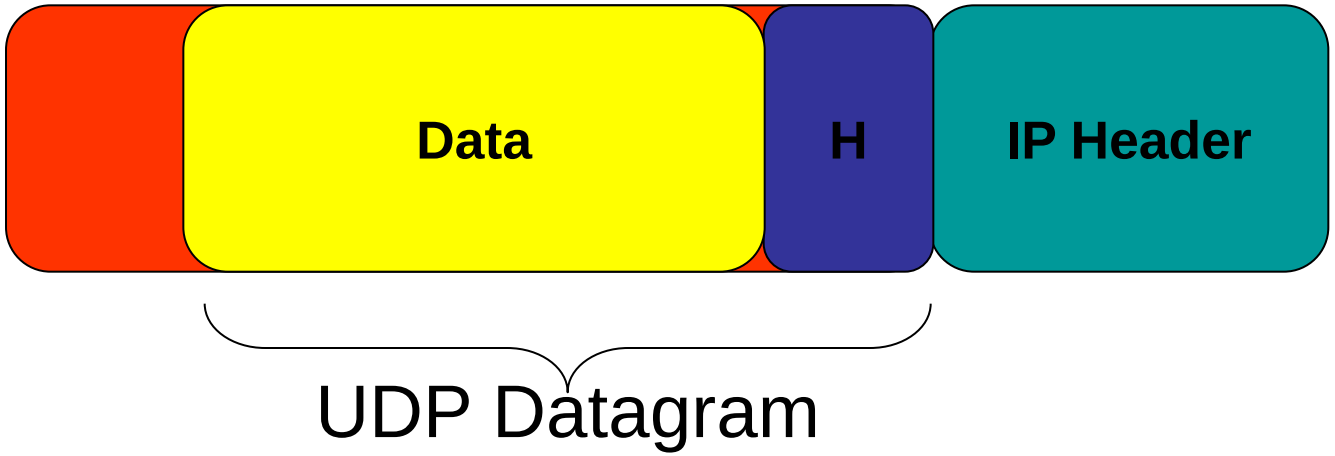

/transport layer/tcp/reliability

- Unit of information sent by TCP is called a *segment*
- Application data is broken into best sized chunks to send
- TCP maintains a timer for receiving an ack from the receiver
- Receiver sends an ack to the sender on receipt of data
- A checksum is maintained on header as well as data
- A receiving TCP re-sequences the segments if they arrive out of order
- Duplicate segment is discarded
- It also provides flow control

/transport layer/tcp/header

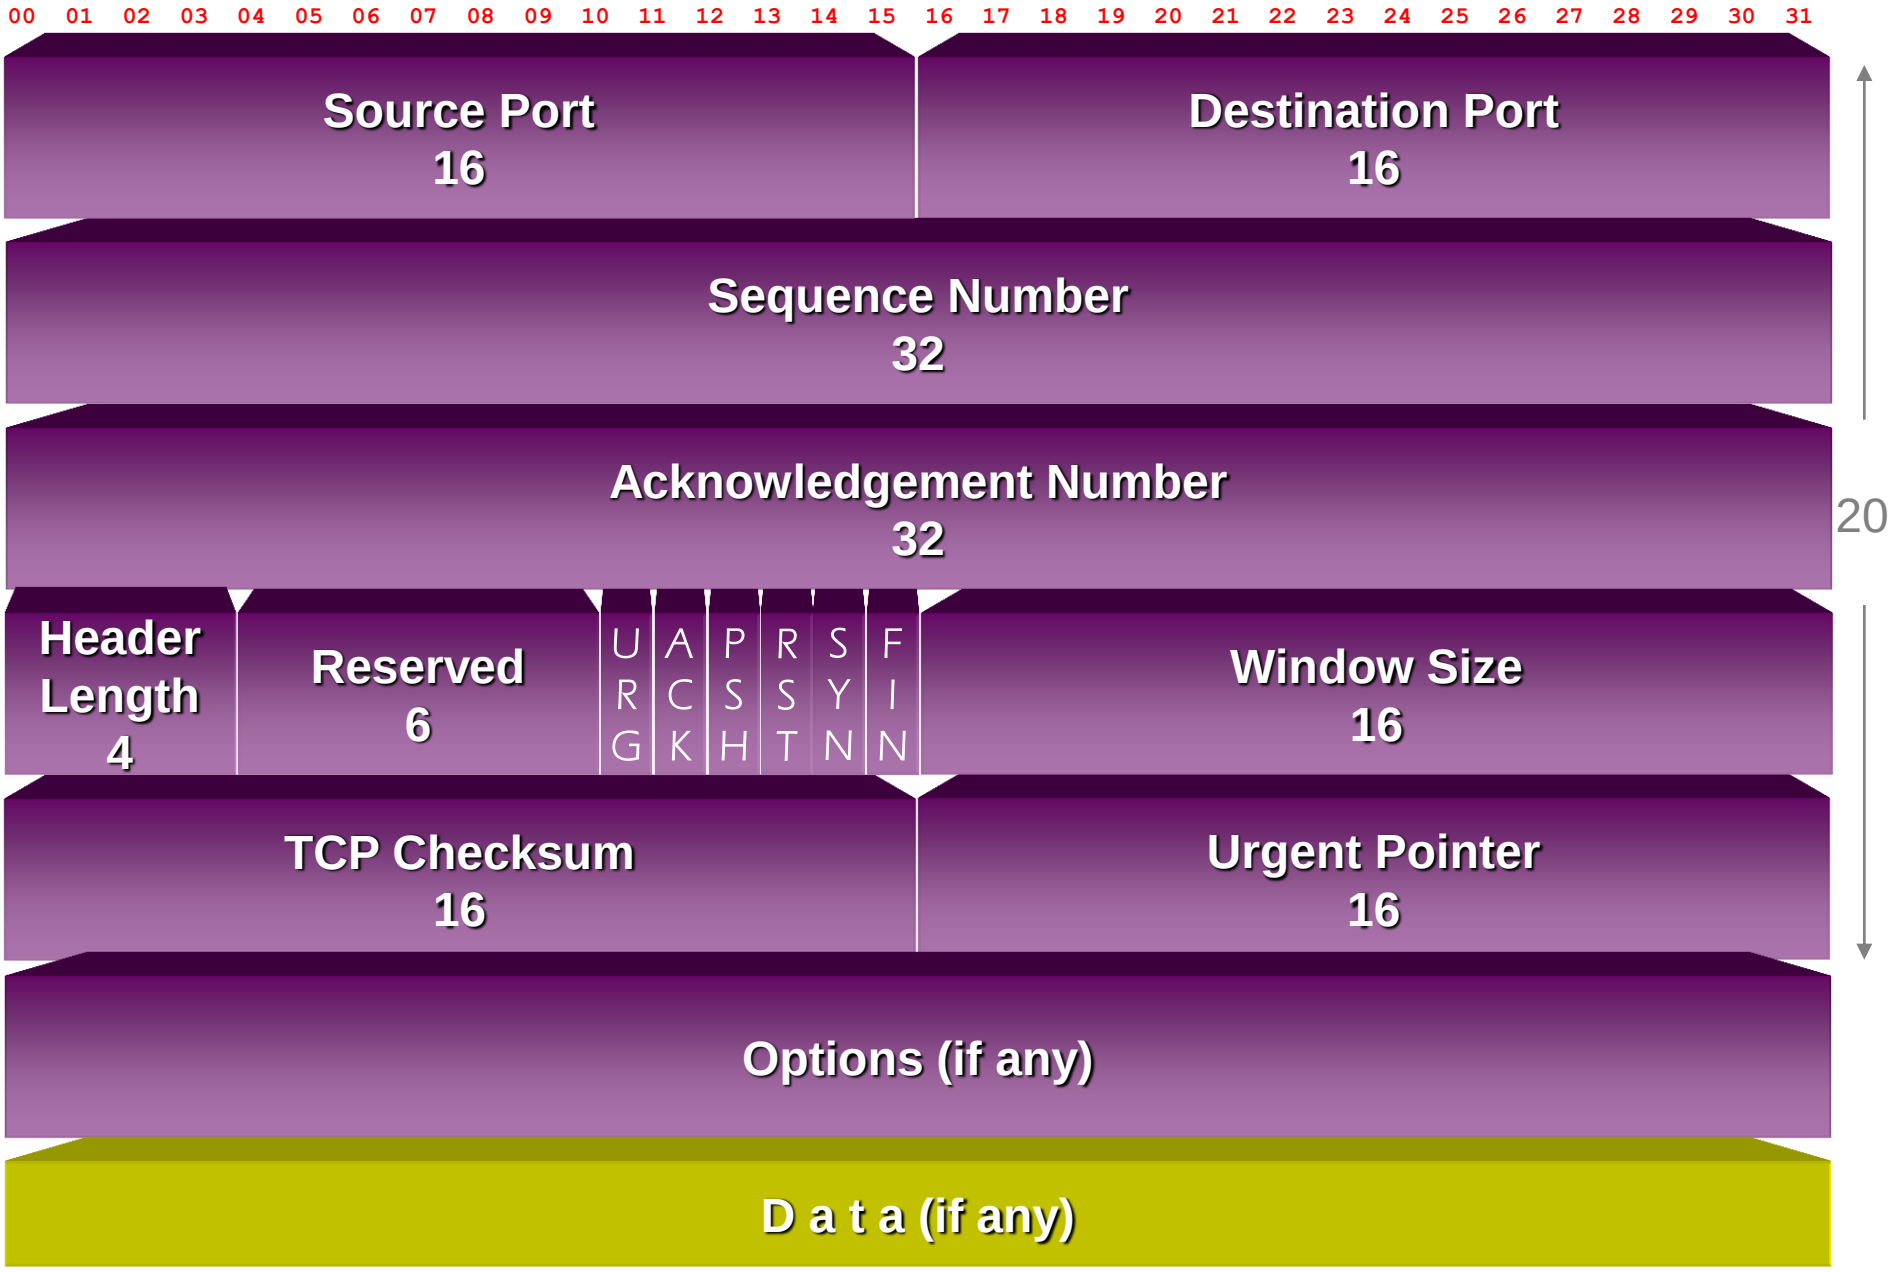

/transport layer/tcp/header/port

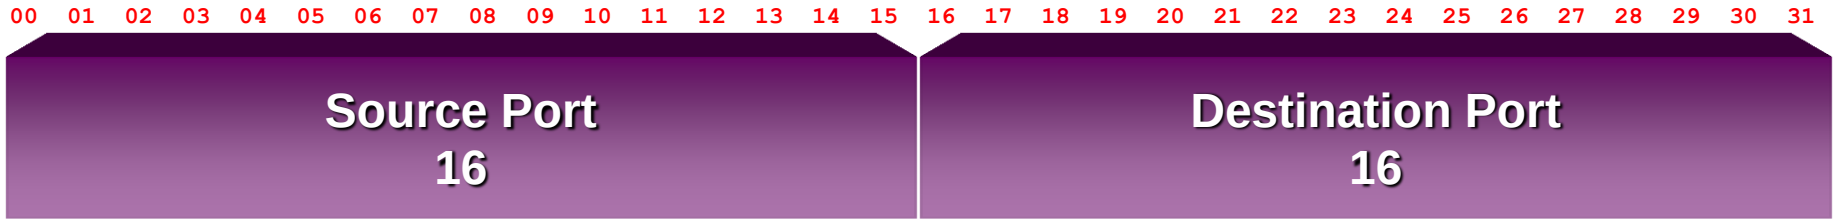

- A port number identifies sending and receiving application
- System port numbers: 1 – 1023 reserved for root (*Managed by IANA*)
- User port numbers 1024 – 49151 are available for user applications (Can be registered with IANA)
- 49152 – 65535 Private/Dynamic port numbers

/transport layer/tcp/header/port

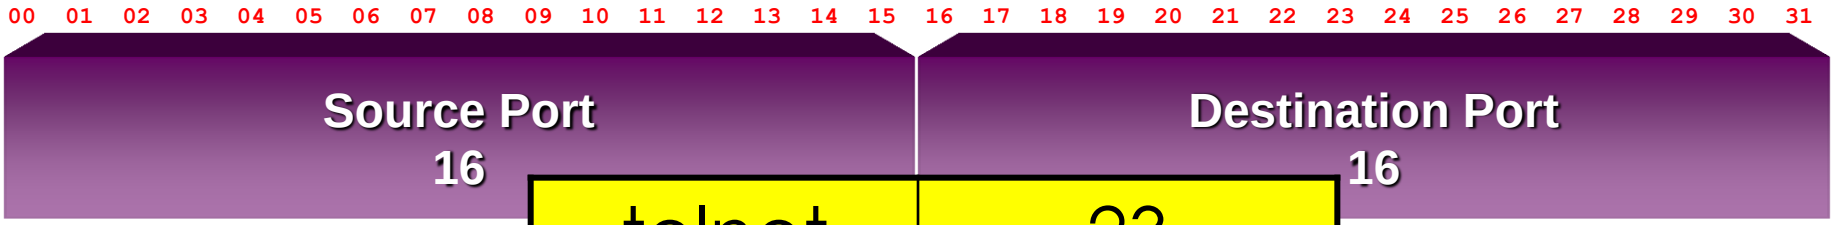

telnet	23
ftp	21
smtp	25
tftp	69

- An IP addre
- Port numbe ion
- The src/ds nd src/dst IP addresses ntify each connectior
- *Socket:* cc P addr, port number
- Well known applications and their port numbers are kept in /etc/services

/transport layer/tcp/header/sequence number

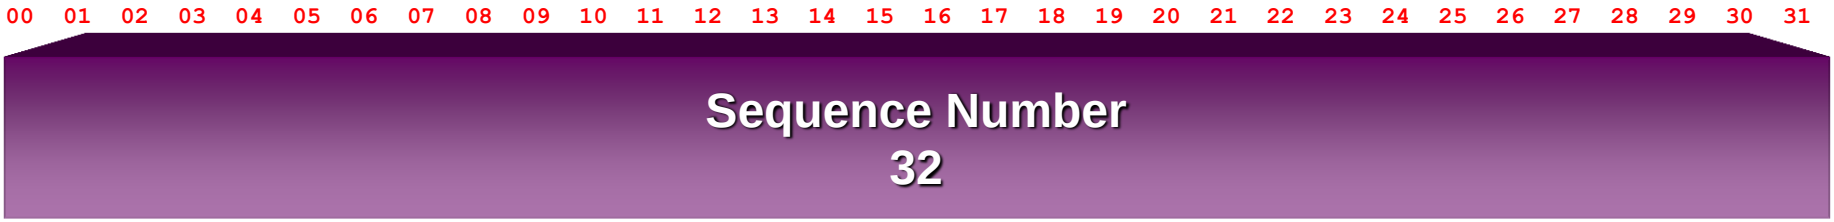

- TCP numbers each byte with a sequence number
- "Sequence Number" field specifies the first byte number from the data being sent to the receiving TCP
- A random *initial sequence number* (ISN) is chosen by the host when a new connection is established
- Sequence number wraps back to 0 after reaching $2^{32} - 1$

/transport layer/tcp/header/acknowledgement number

- Used by receiver to acknowledge the number of bytes received so far by specifying the next byte number it expects
- This field is valid only if ACK flag is set
- Each end of connection maintains sequence number of the data flowing in each direction

/transport layer/tcp/header/length

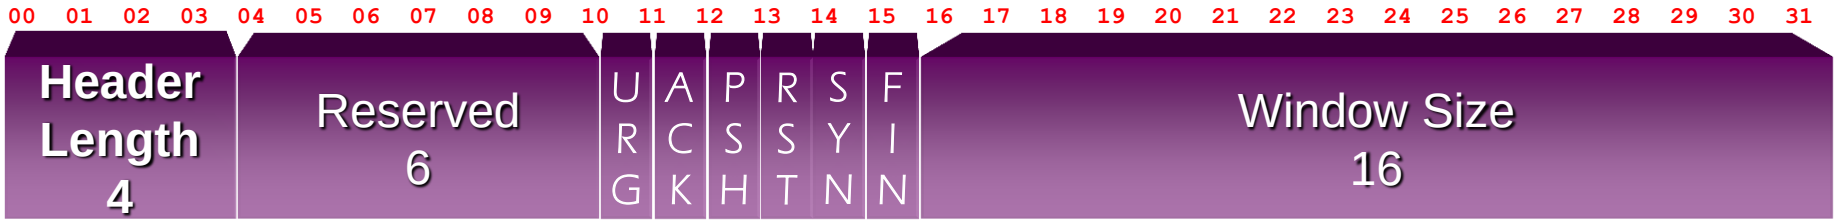

- Header length is in 32-bit words
- Length of options field is variable
- Max length = 60, Min = 20 bytes
- Header always has to be in multiple of 32-bit words:
  - If its not, options have to be padded

/transport layer/tcp/header/flags

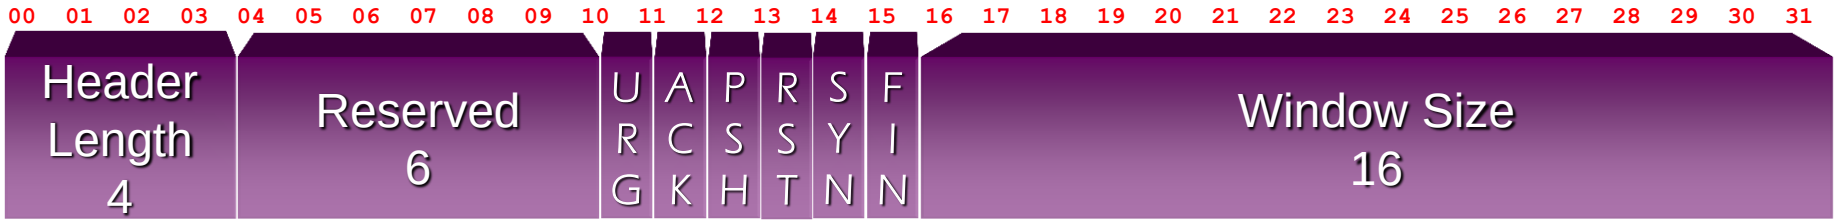

- URG – *The urgent pointer is valid (usage: interrupting data transfer)*
- ACK – *Acknowledgement number is valid*
- PSH – *The receiver should pass this data to the application as soon as possible (eg. telnet)*
- RST – *Reset the connection*
- SYN – *Synchronize sequence number to initiate connection*
- FIN – *The sender is finished sending data*

/transport layer/tcp/header/window size

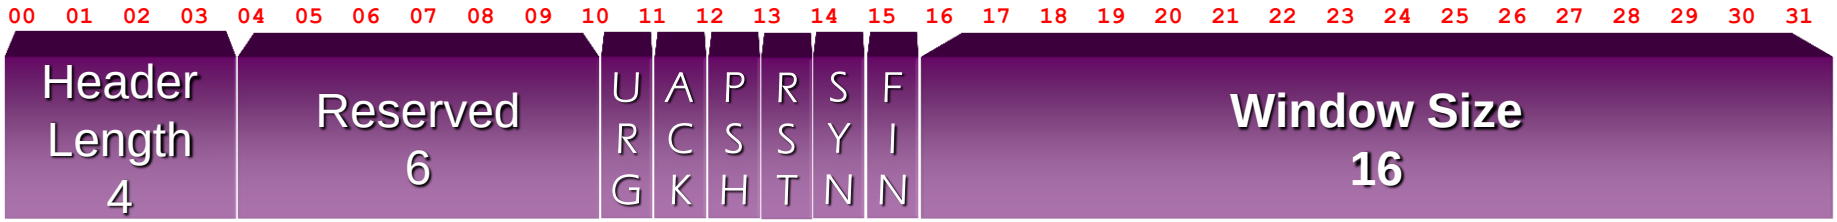

- TCP uses this window to advertise how many bytes it can receive at a time
- Useful for flow control
- Sliding window protocol based on this

/transport layer/tcp/header/checksum

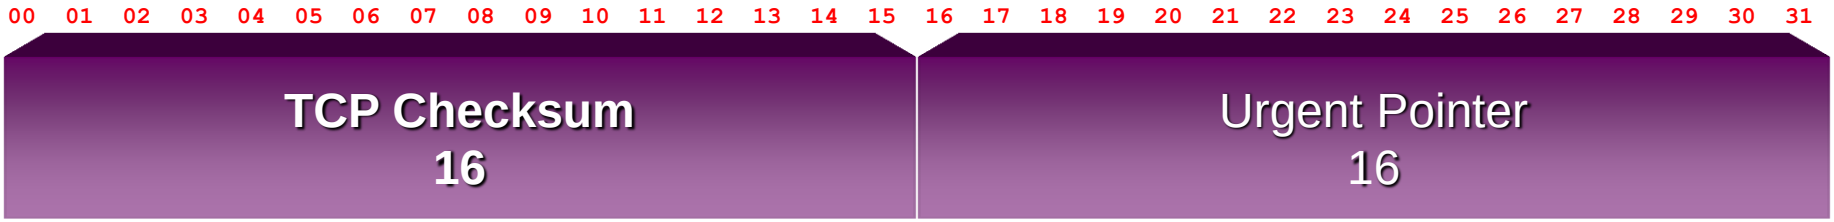

- Mandatory field
- Covers header and data
- Uses a pseudo-header for calculation:
  - IP Source and Destination Address fields
  - IP Protocol field
  - TCP Length field

/transport layer/tcp/header/urgent pointer

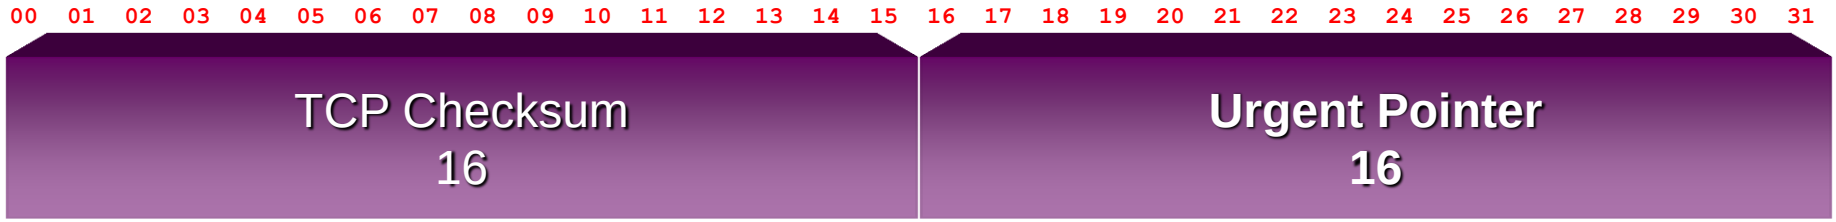

- Urgent pointer is a positive offset that must be added to sequence number, to yield the sequence number of the last byte of urgent data
- Used for transmitting emergency data
- Up to the application as to how the urgent data is to be used
- Urgent pointer is valid only if URG flag is set to 1

/transport layer/tcp/header/options

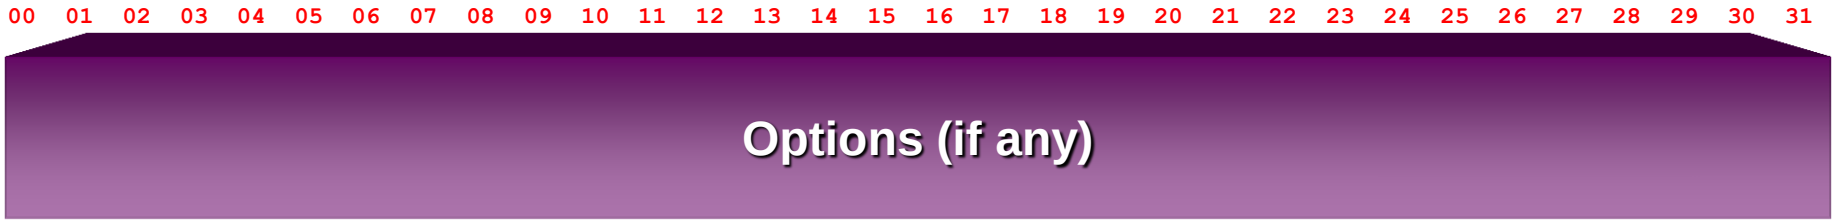

- Maximum segment size
- Window scale factor
- Timestamp
- End of option list
- No operation

Option-Kind (1)	Option-Length (1)	Option-Data (var)

/transport layer/tcp/header/data

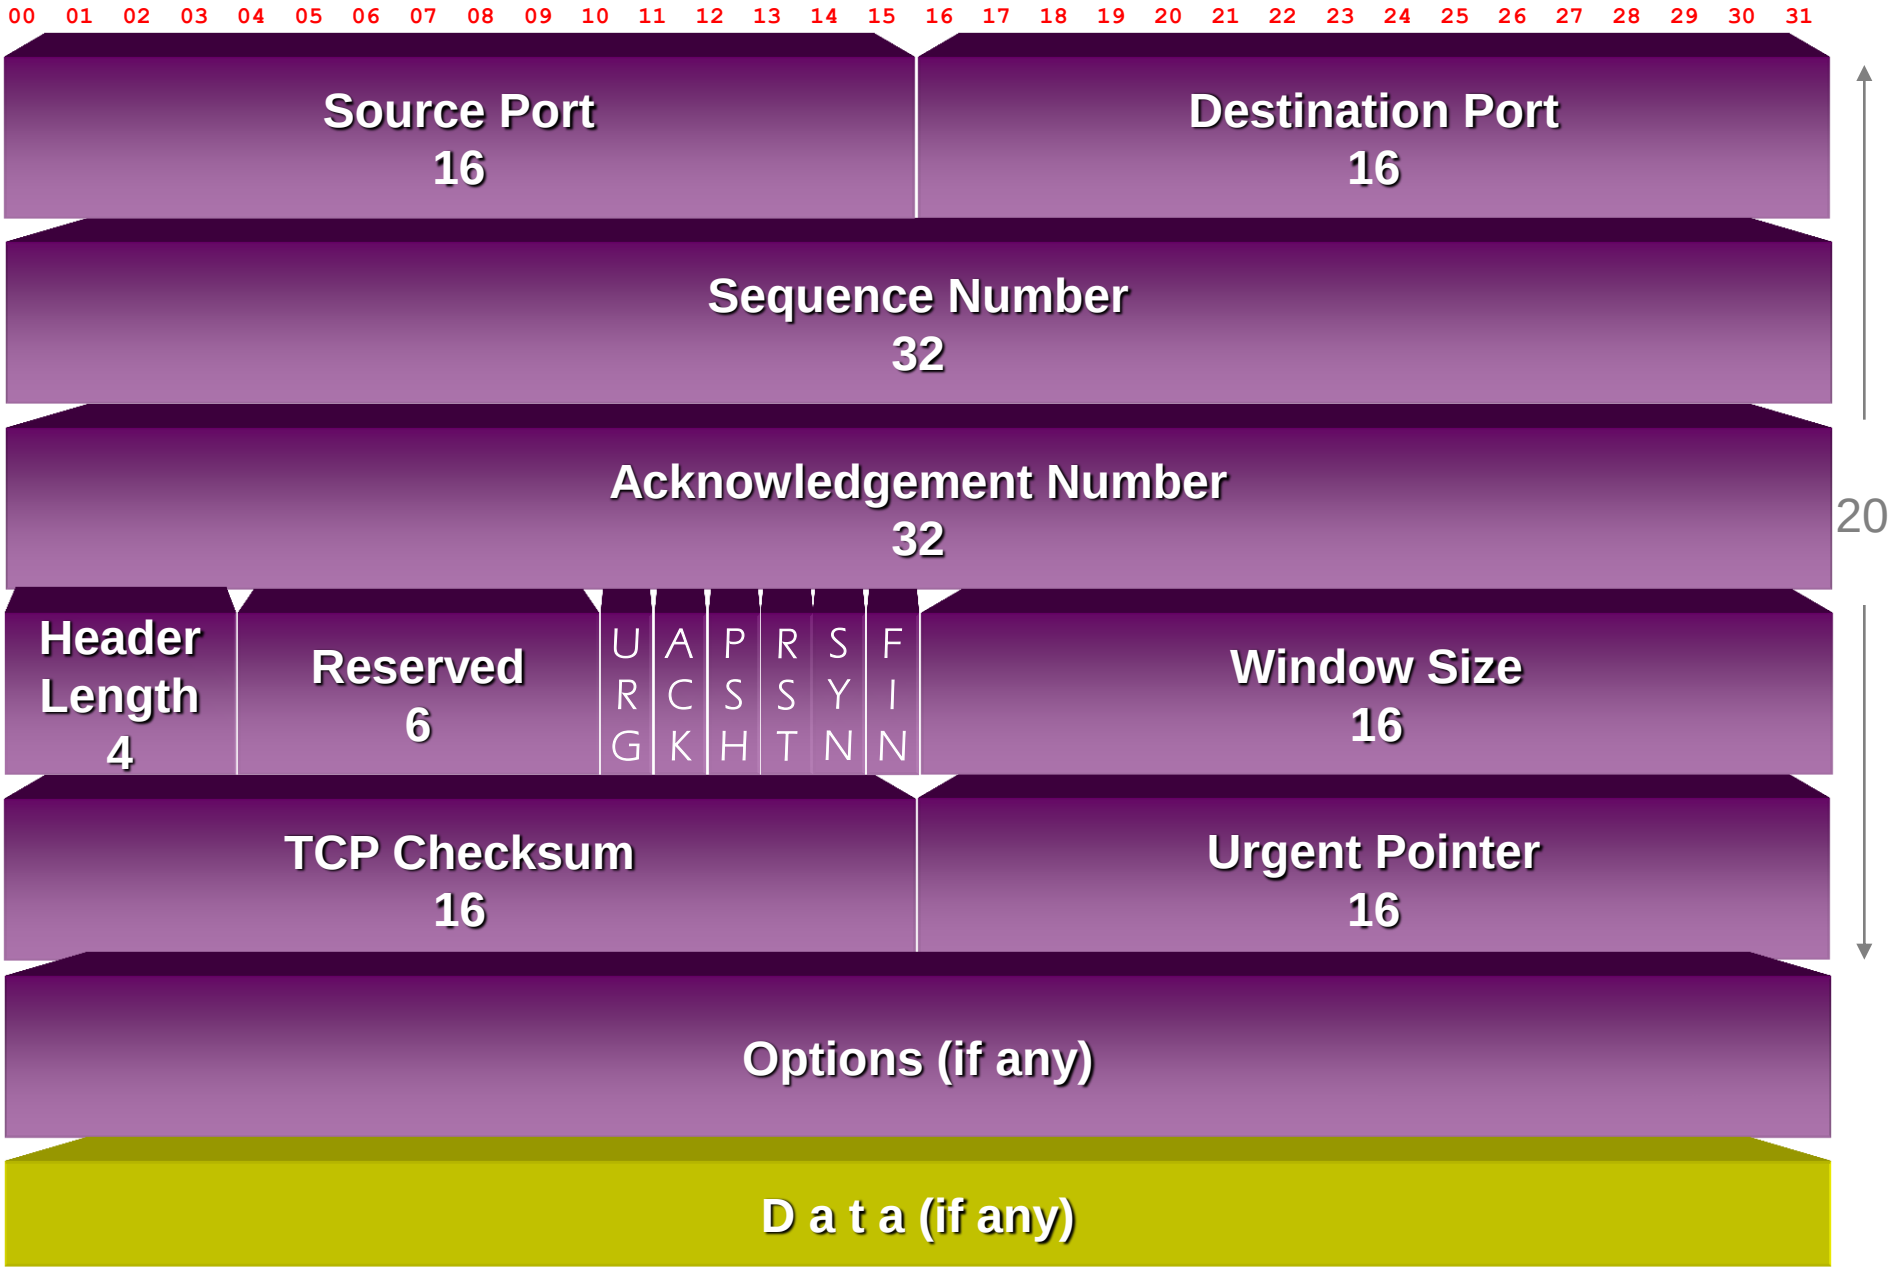

/transport layer/tcp/connection/open

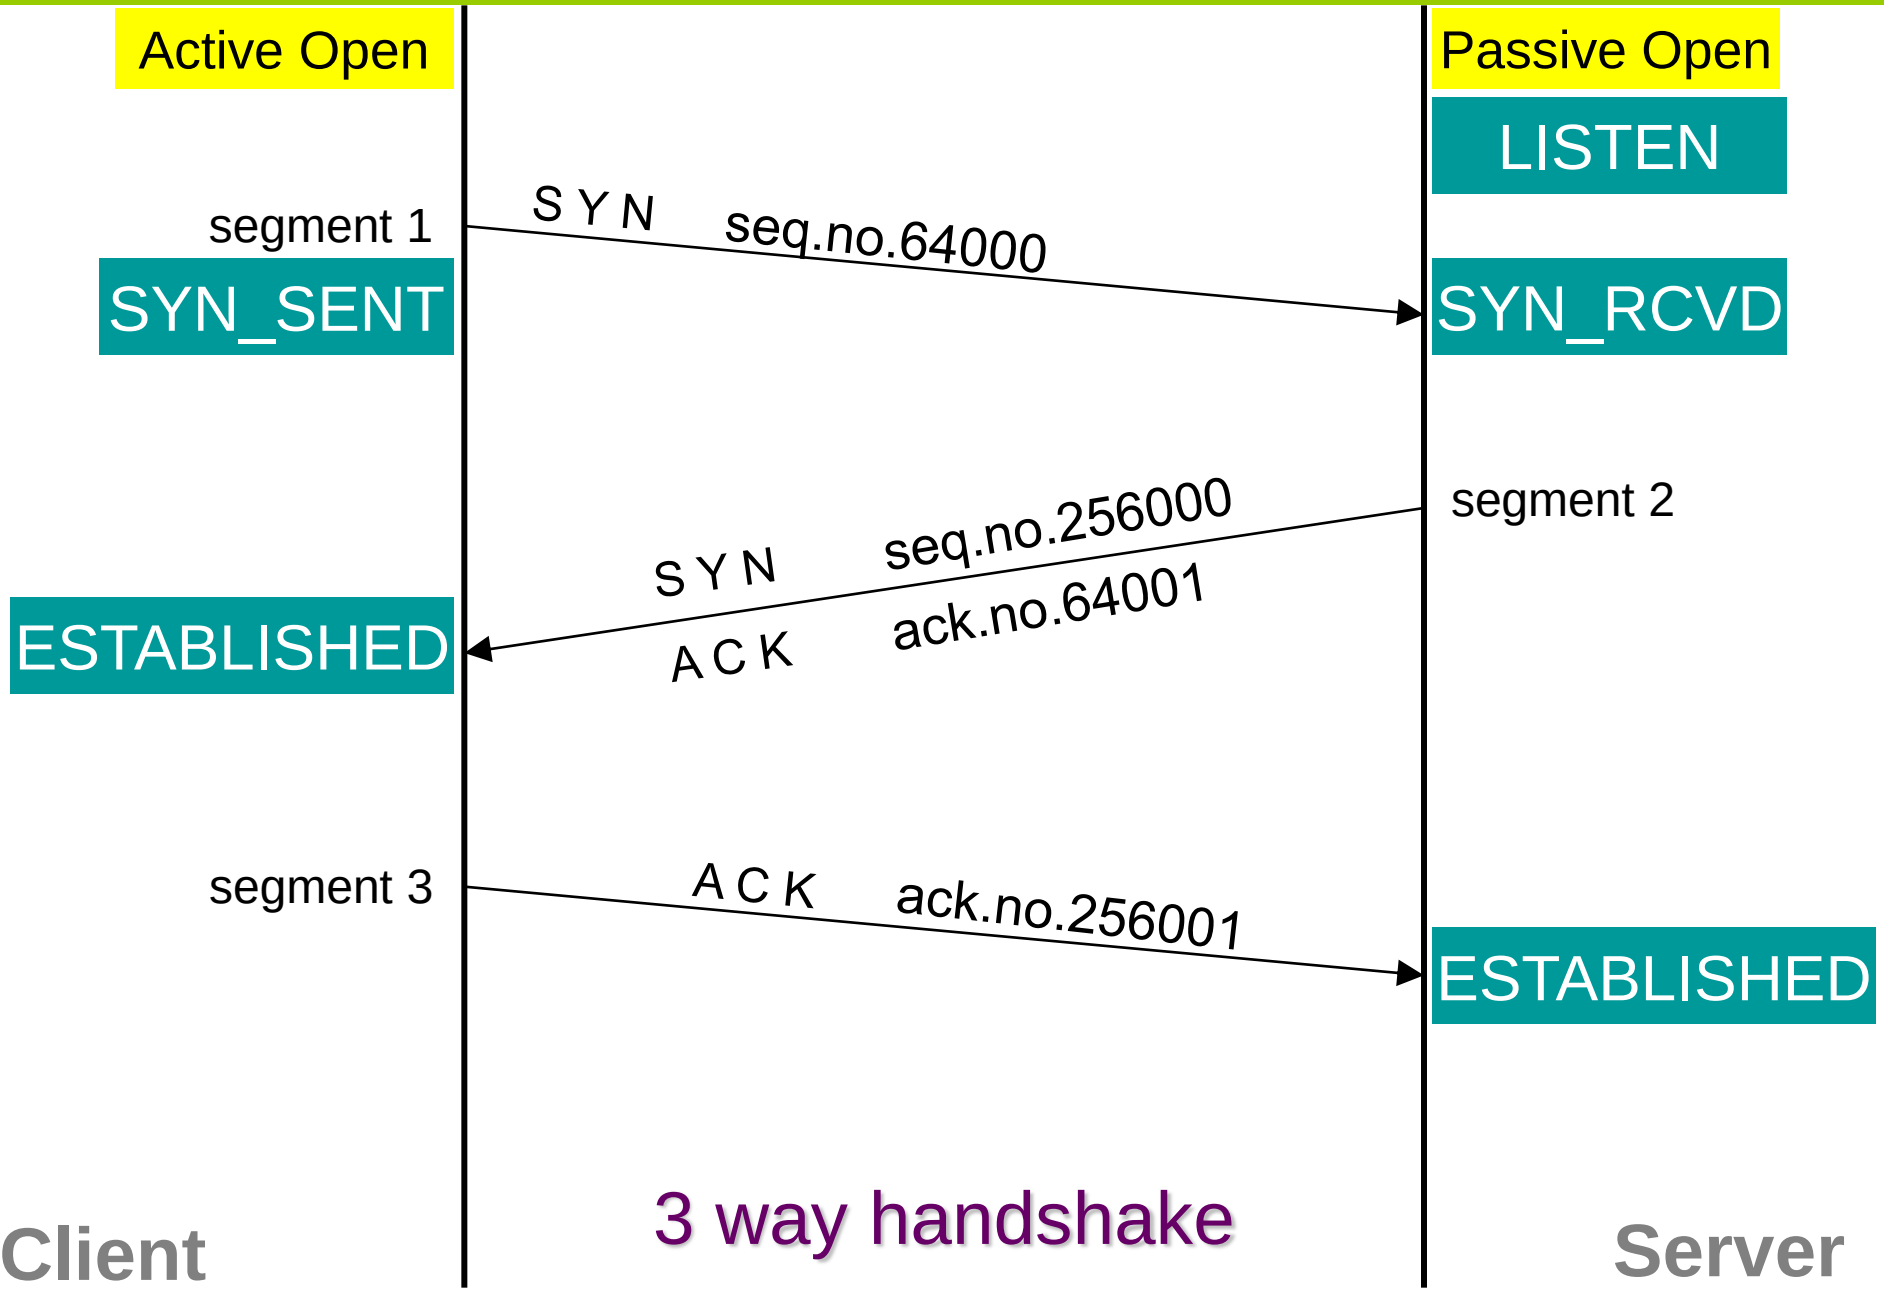

/transport layer/tcp/connection/close

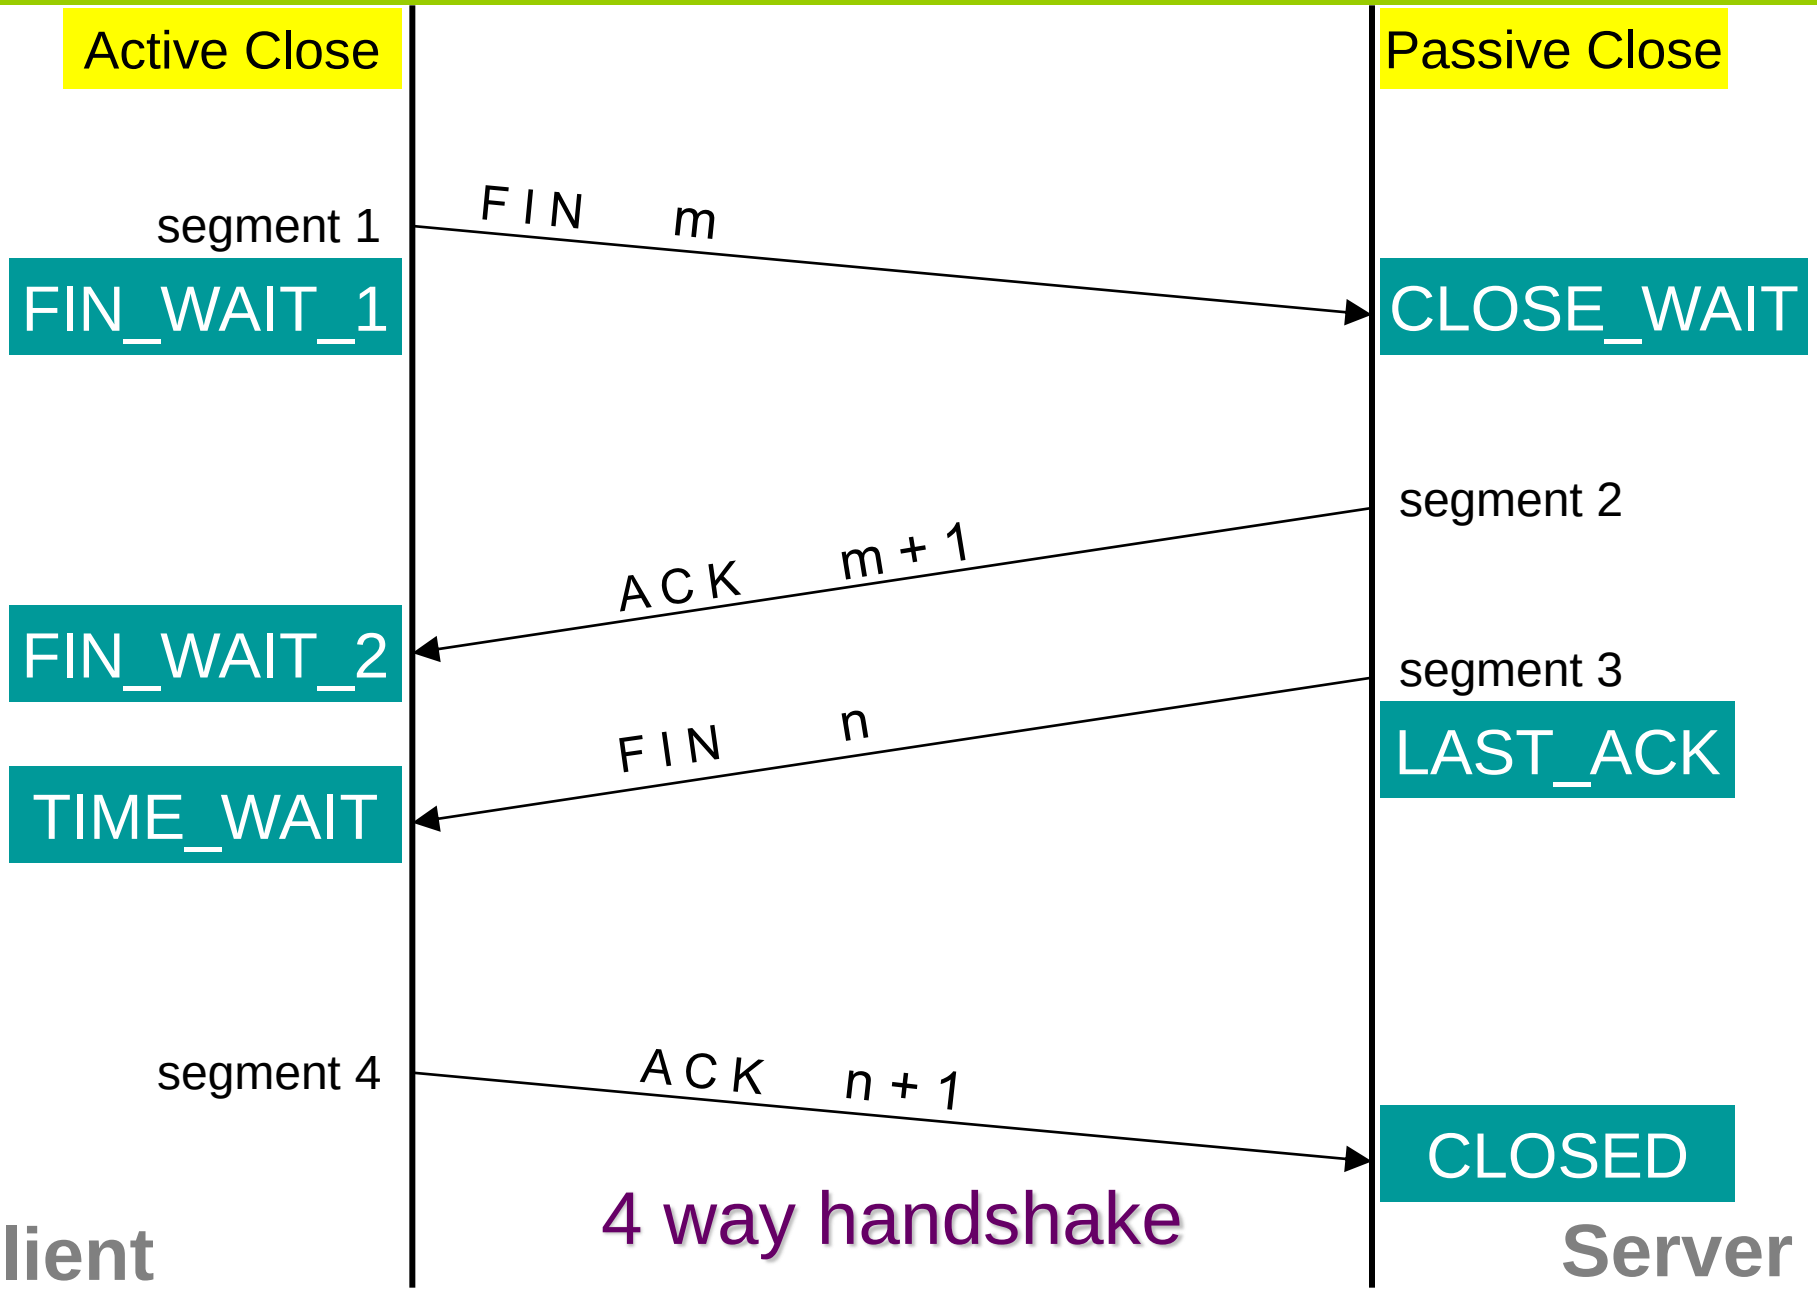

/transport layer/tcp/connection/half close

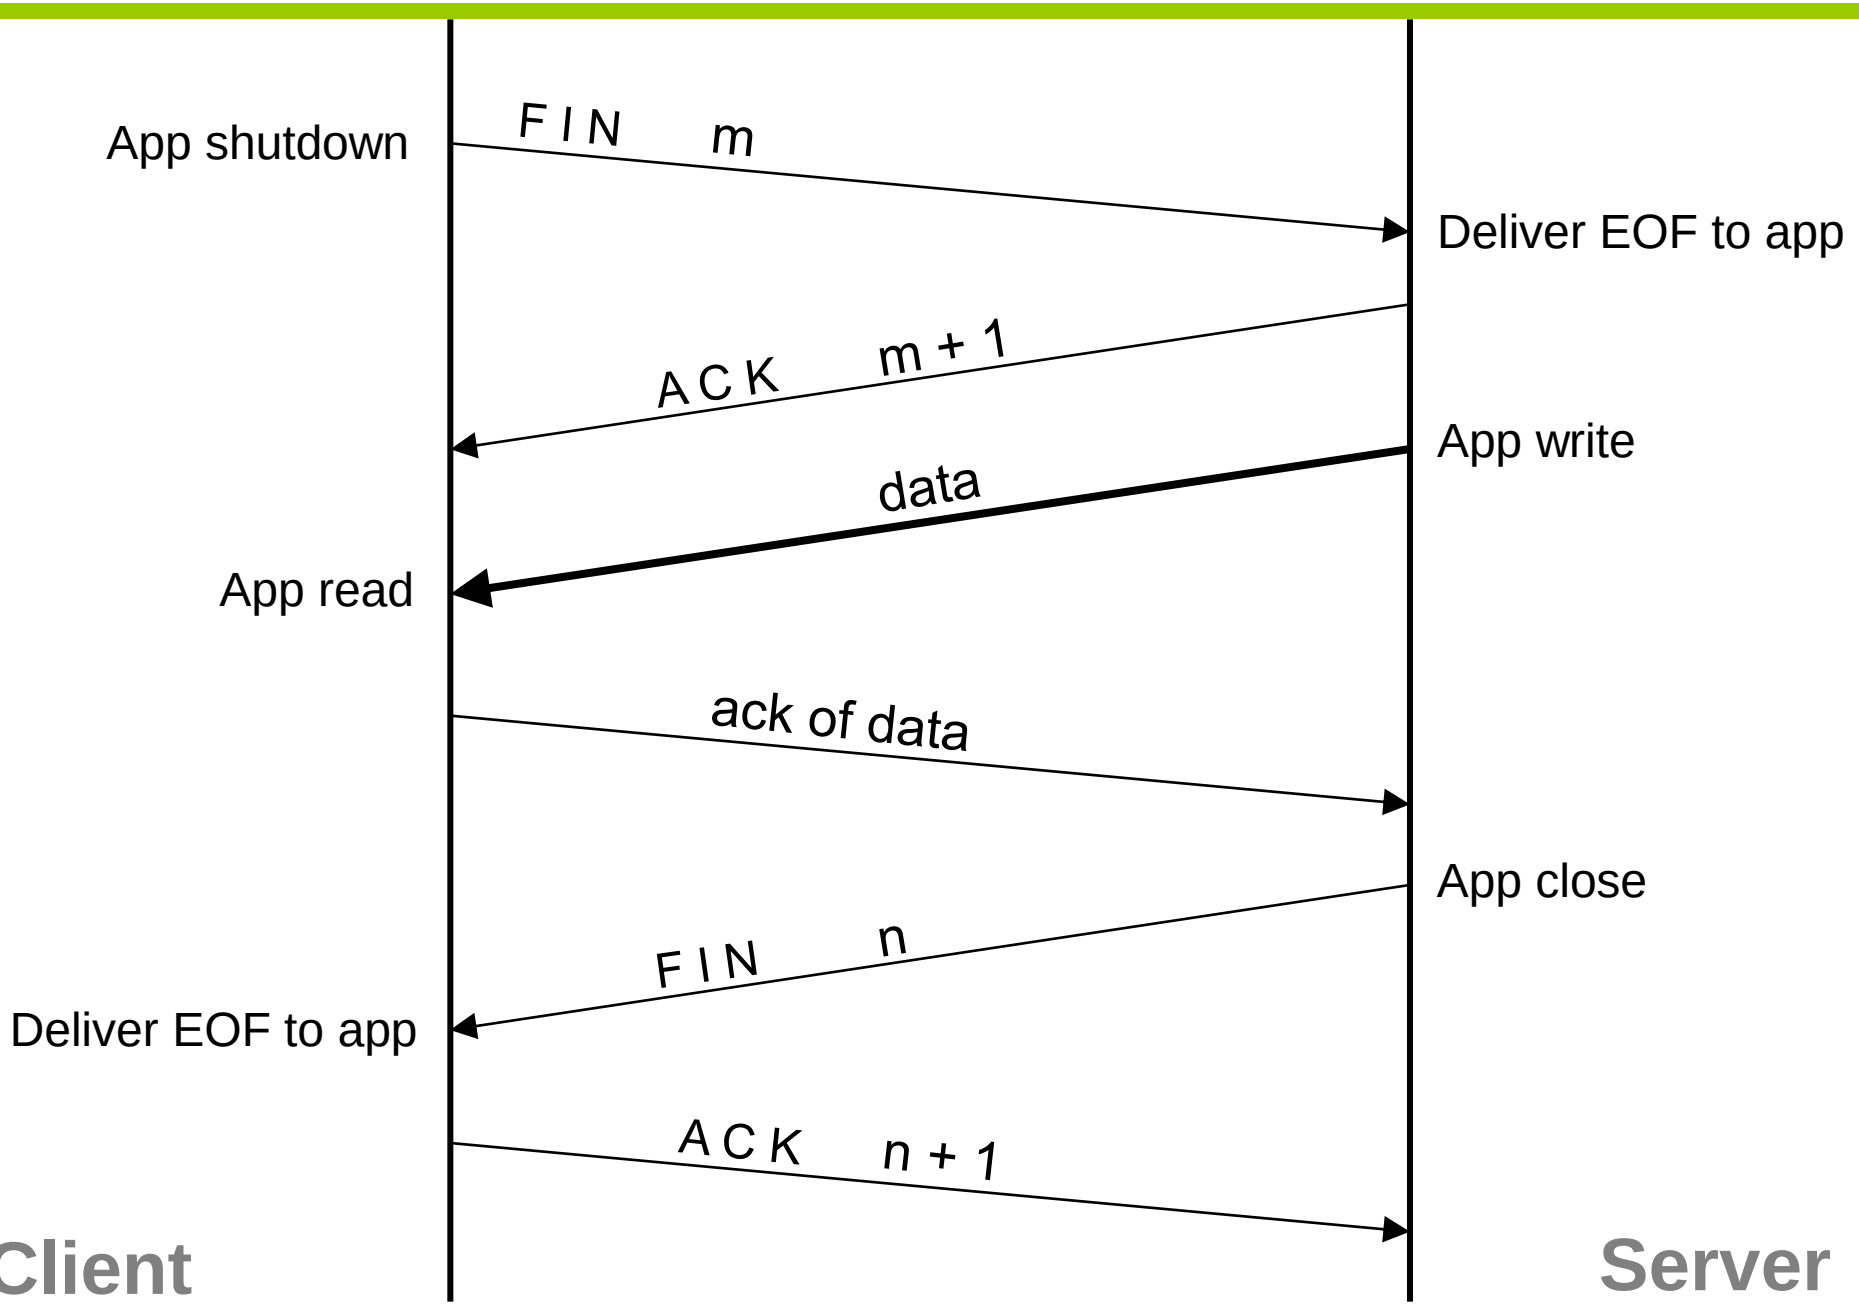

/transport layer/tcp/connection/simultaneous

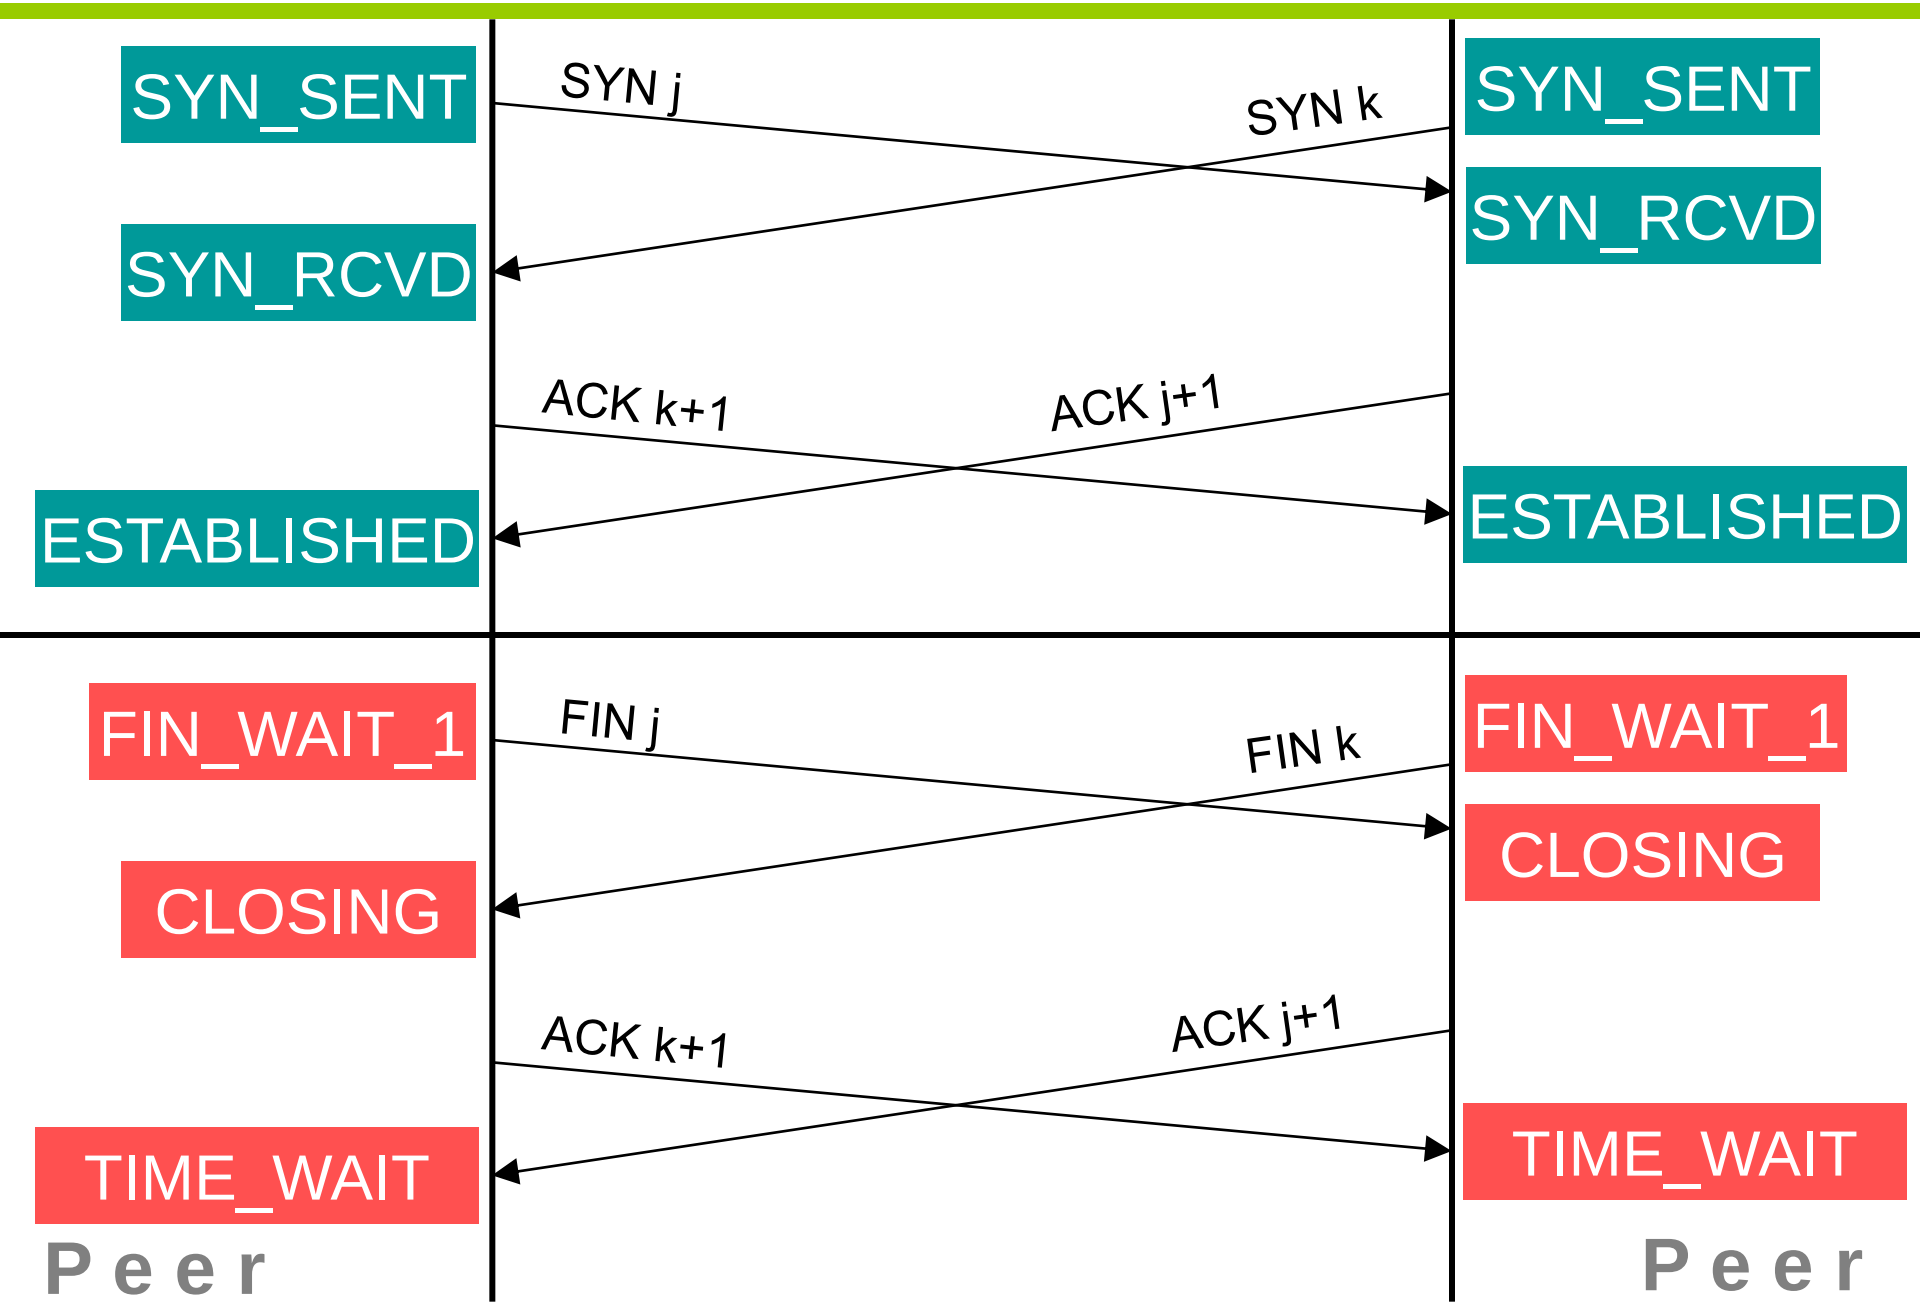

- Connection can't be established if remote host is down
- There is a timeout set for a retry
- Maximum segment lifetime (MSL): *It is the max amount of time any segment can exist in the network before being discarded*
- A connection is reset:
  - If request arrives and no process is listening on the destination port
  - By sending a reset (*abortive release*)
- A TCP connection is said to be *half-open* if one end has closed or aborted the connection without the knowledge of the other end

/transport layer/tcp/connection/queue

- There's a fixed length queue of connection
- Backlog is between 0 – 5
- Connections in a queue are already accepted by TCP, they are waiting to be accepted by the application
- If there is no room in the queue, TCP simply ignores the SYN from incoming connection

- Interactive (*E.g. rlogin, telnet*)
  - Normally transmitted in segments smaller than the max segment size
  - Delayed acks are piggybacked by receiver along with data going back to sender over WAN
- Bulk
  - Sliding window protocol: Receiver does not have to acknowledge every received packet
  - The acks are cumulative
  - Window is advertised by receiver

- PUSH Flag
- Slow start
- URGENT Flag
- Congestion

1. Trap TCP packets using *tcpdump*
2. Trap TCP packets belonging to specific application (telnet, ftp, etc.) between any two hosts
3. Use *netstat* to find how many TCP sockets are open and what their states are

/transport layer/udp/overview

- Created as an alternative transport protocol for applications that don't need the features of TCP
- Proposed in RFC 768 in 1980
- Serves as an interface between application processes and IP
- Simple and fast

/transport layer/udp/header

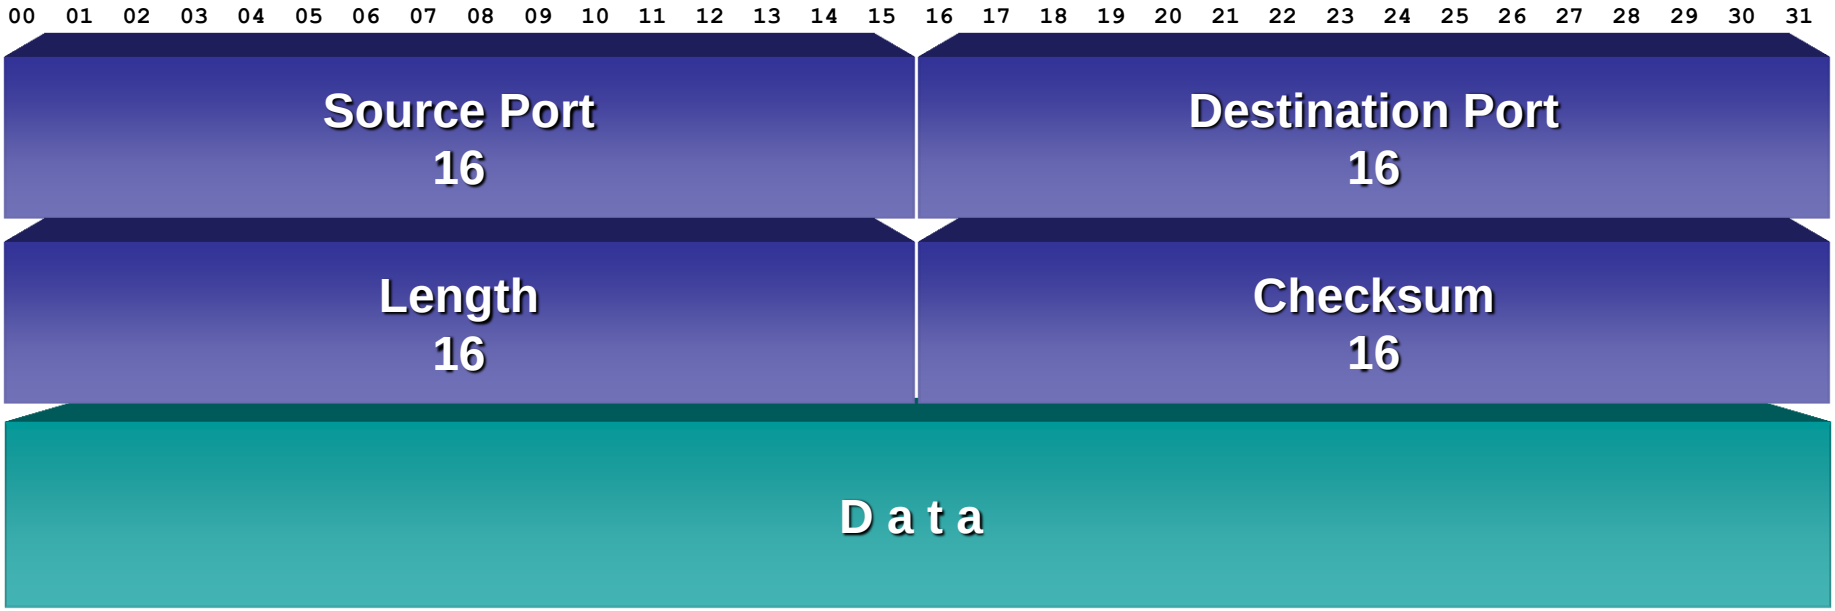

- Checksum field is optional
- Checksum is computed for actual header + pseudo header comprising of:
  - IP Source and Destination Address fields
  - IP Protocol field
  - UDP Length field

/transport layer/udp/does not

- Establish connections before sending data. It packages the data and sends it off

- Provide acks
- Provide guarantee of reception

- Detect lost messages and retransmit them

- Ensure data ordering
- Provide any mechanism to handle congestion or manage the flow of data between devices

/transport layer/udp/applications

Port#	Keyword	Protocol	Comments
53	domain	Domain Name Server	Uses a simple request / reply messaging system
67 / 68	bootps / bootpc	Bootstrap protocol & DHCP	Host configuration protocols
69	tftp	Trivial File Transfer Protocol	For quick and easy transfer of small files
161 / 162	snmp	Simple Network Management Protocol	An administrative protocol
520 / 521	router / ripng	Routing Information Protocol (RIP-1, RIP-2, RIPng)	Routing protocols
2049	nfs	Network File System	Used UDP earlier. New versions use TCP

1. Trap UDP datagrams using *tcpdump*

2. Trap UDP datagrams and dump the contents for analysis of its header

Network Layer

/network layer/ip

- This layer is responsible for routing messages through networks
- IP is a connectionless protocol that doesn't provide reliability, flow control or error recovery (These functions must be provided at a higher level)
- It offers a best effort service. If something goes wrong, IP discards the datagram and tries to send an ICMP message to the source host
- A message unit in an IP network is called an IP *datagram*. This is the basic unit of information transmitted across TCP/IP networks

/network layer/ip

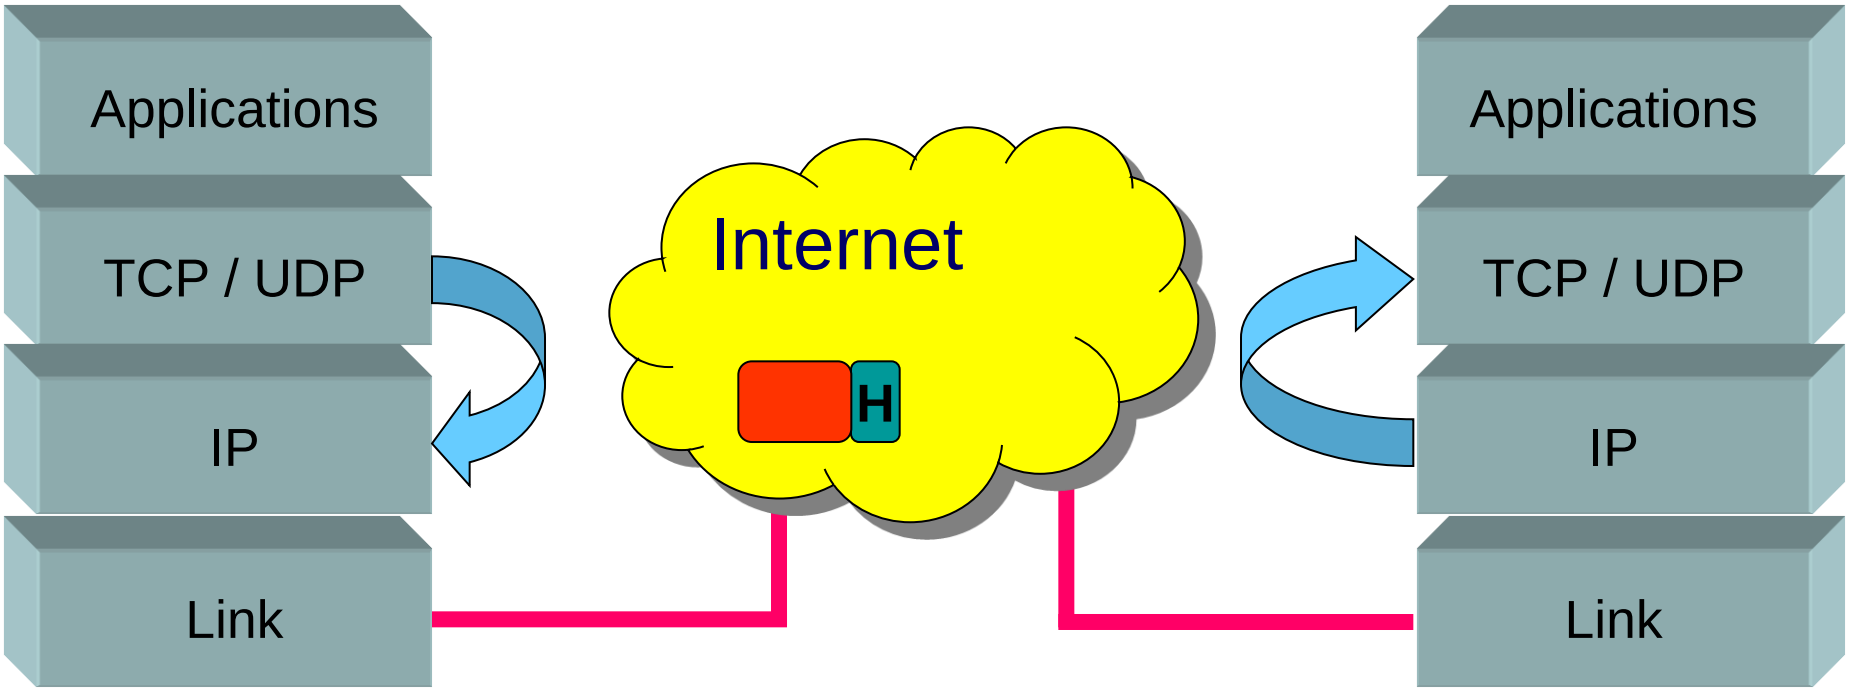

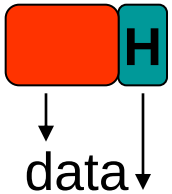

IP Datagram

header – Contains control and addressing information

/network layer/ip/devices

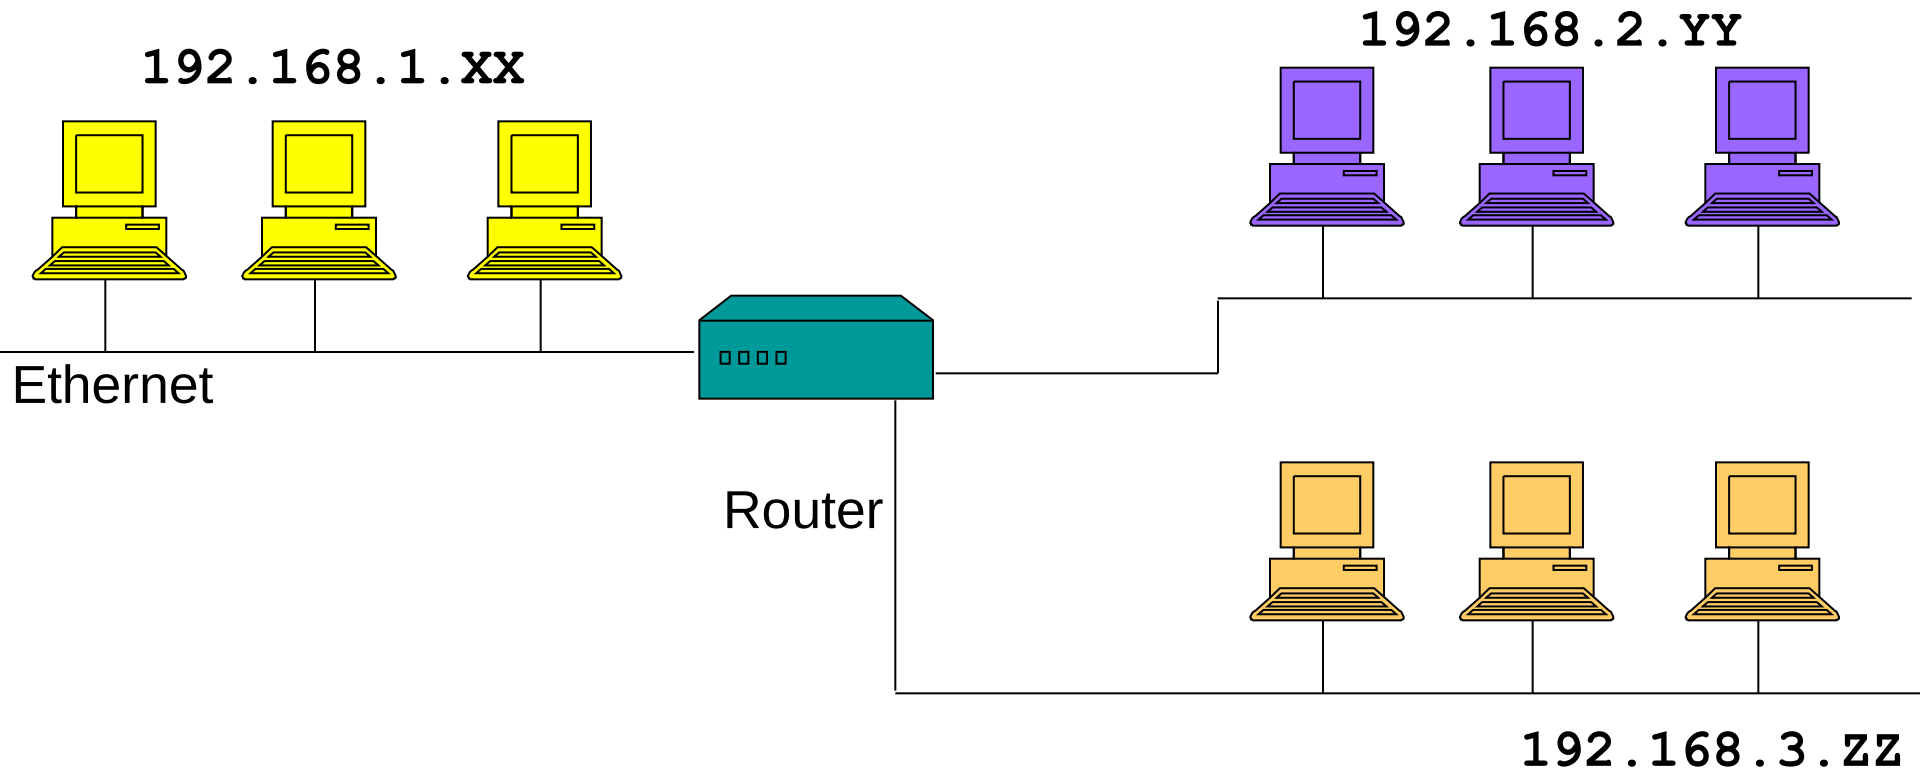

- A router provides interface between two networks. Also called as *Gateway*
- It routes the *datagrams* leaving and entering the network to enable them to get nearer to their destination

/network layer/ip/devices

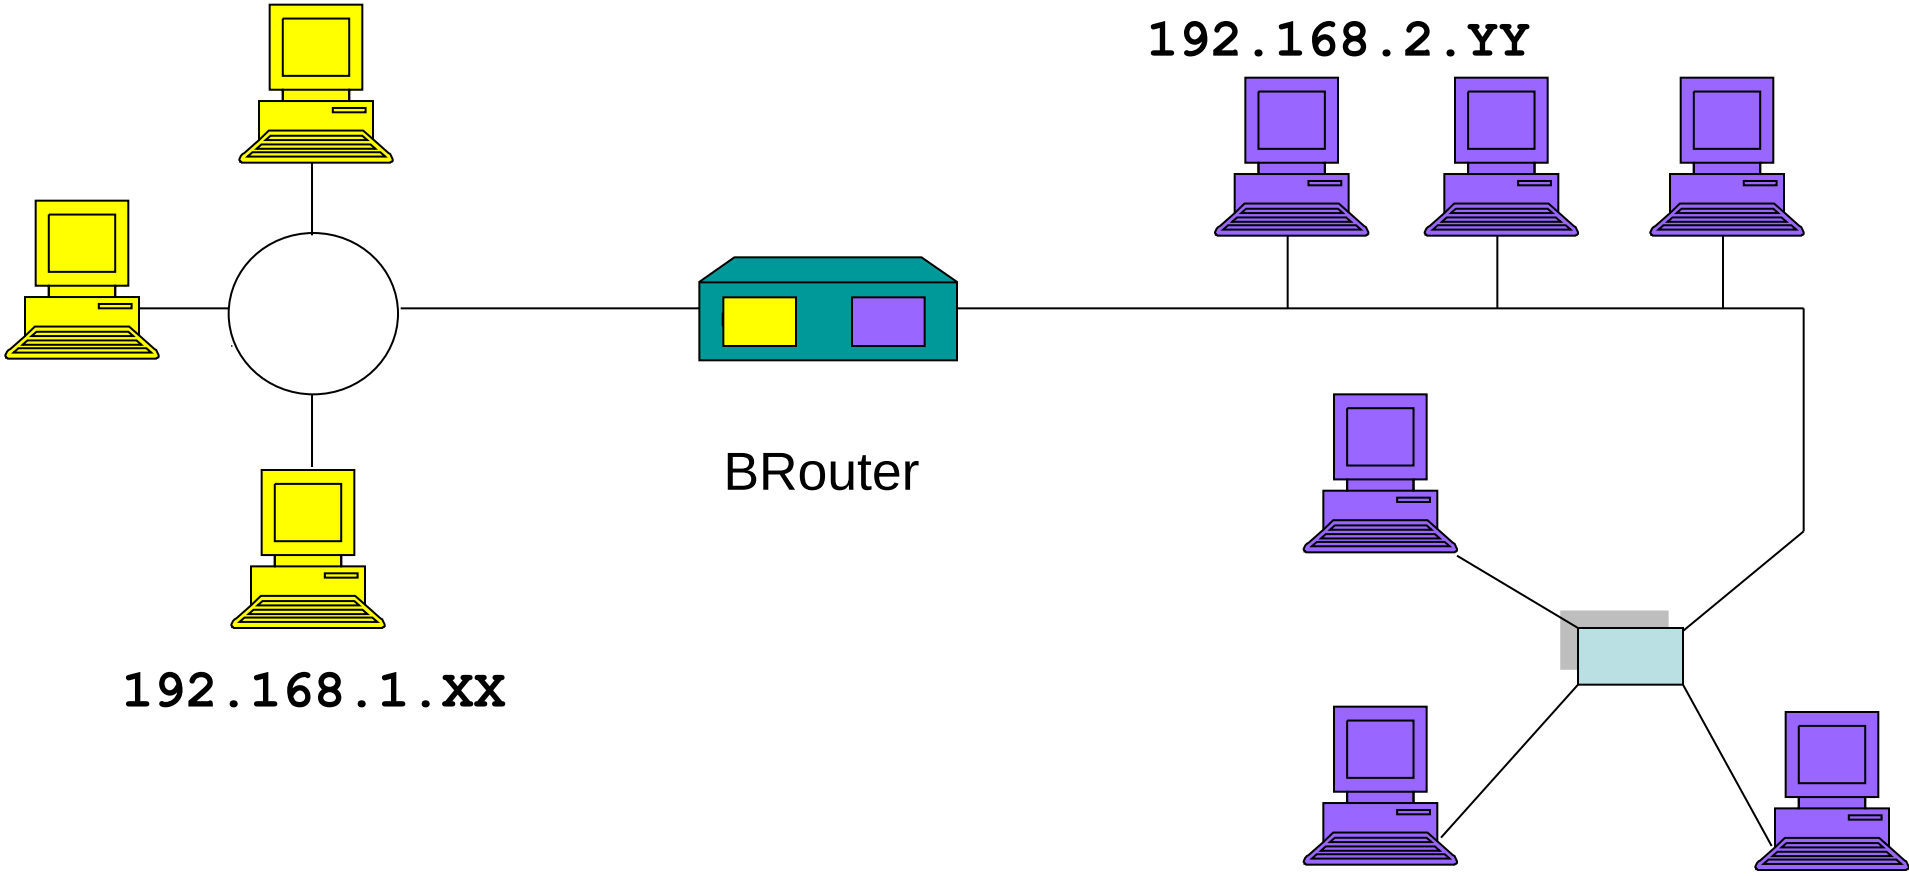

- Wherever necessary the router will translate the network access protocols used by one network into the protocols used by the other

/network layer/ip/header

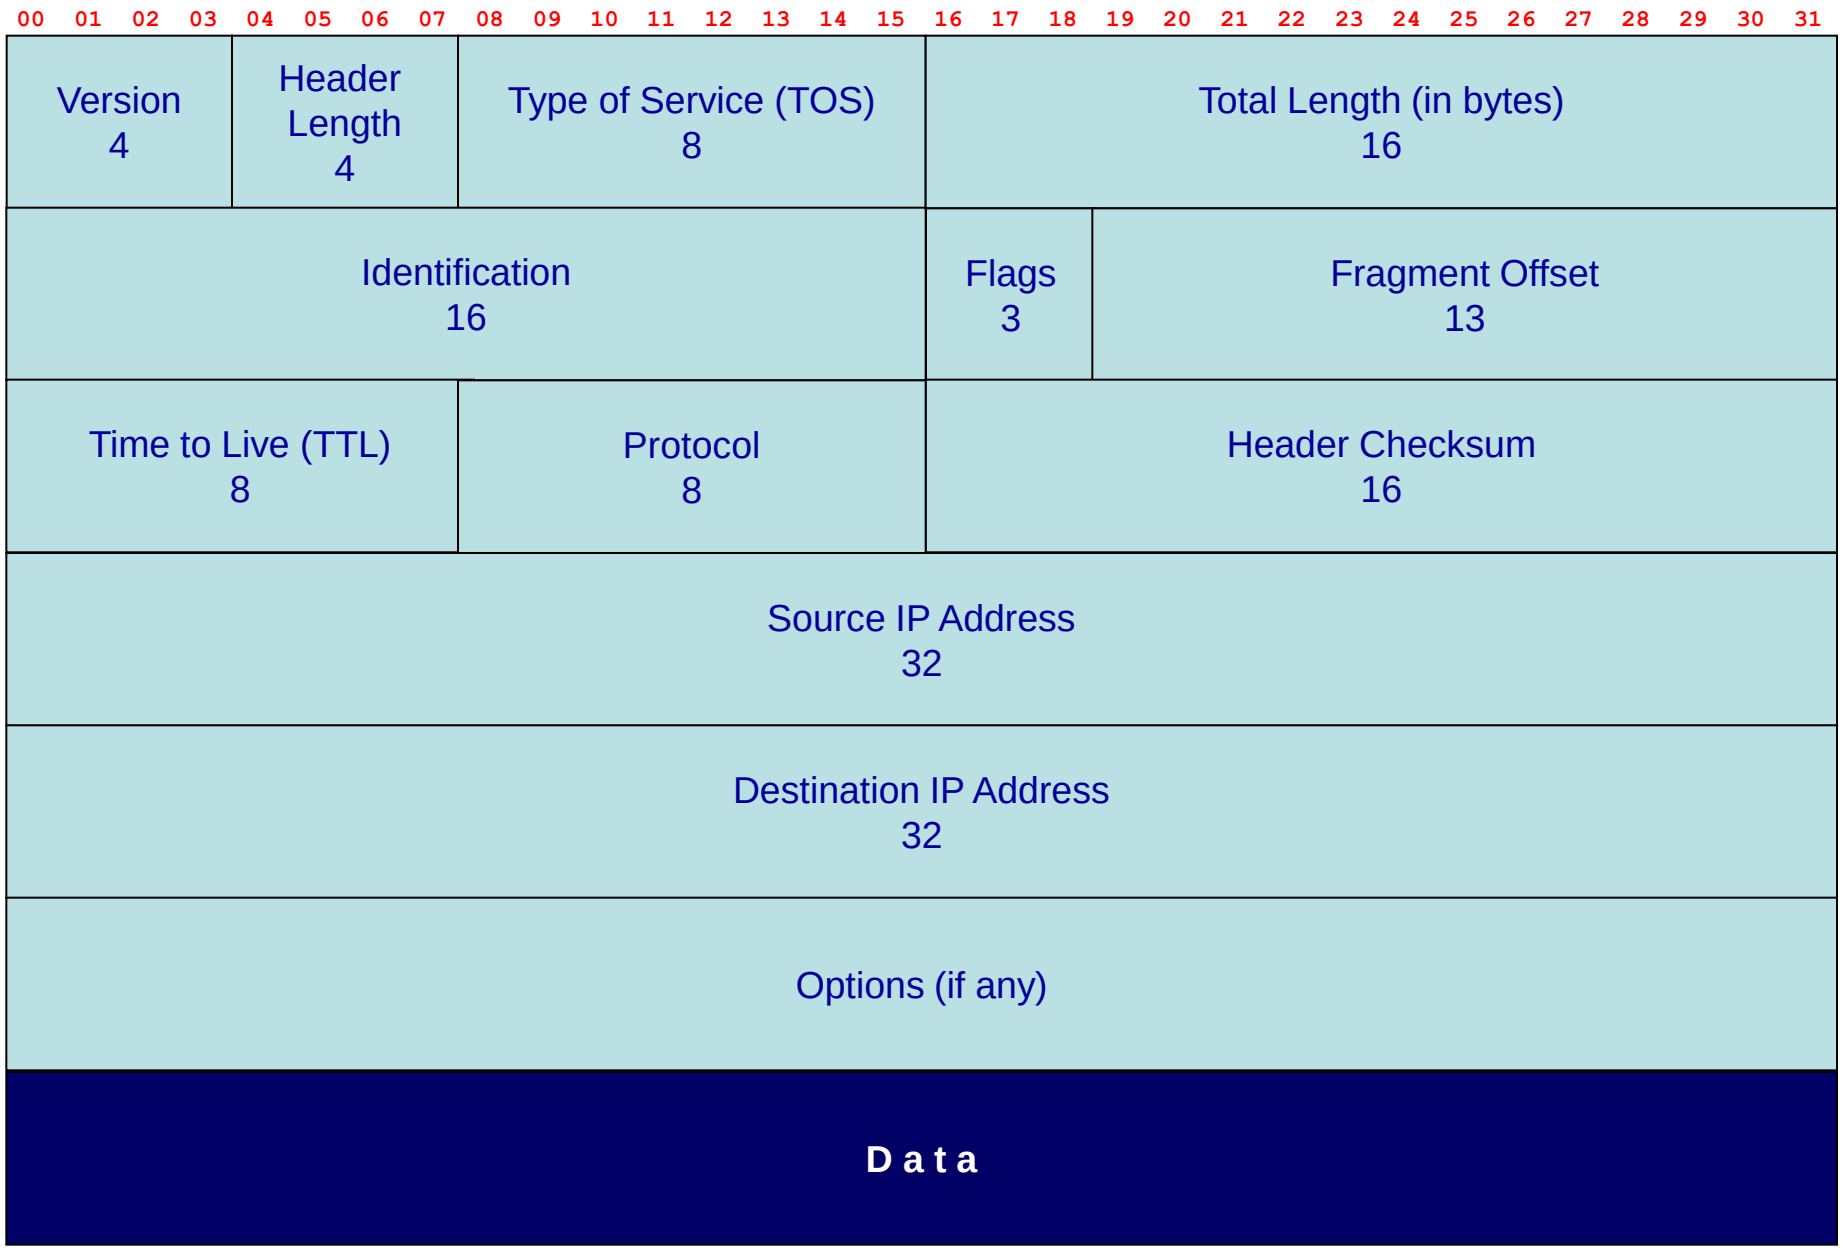

/network layer/ip/header/version

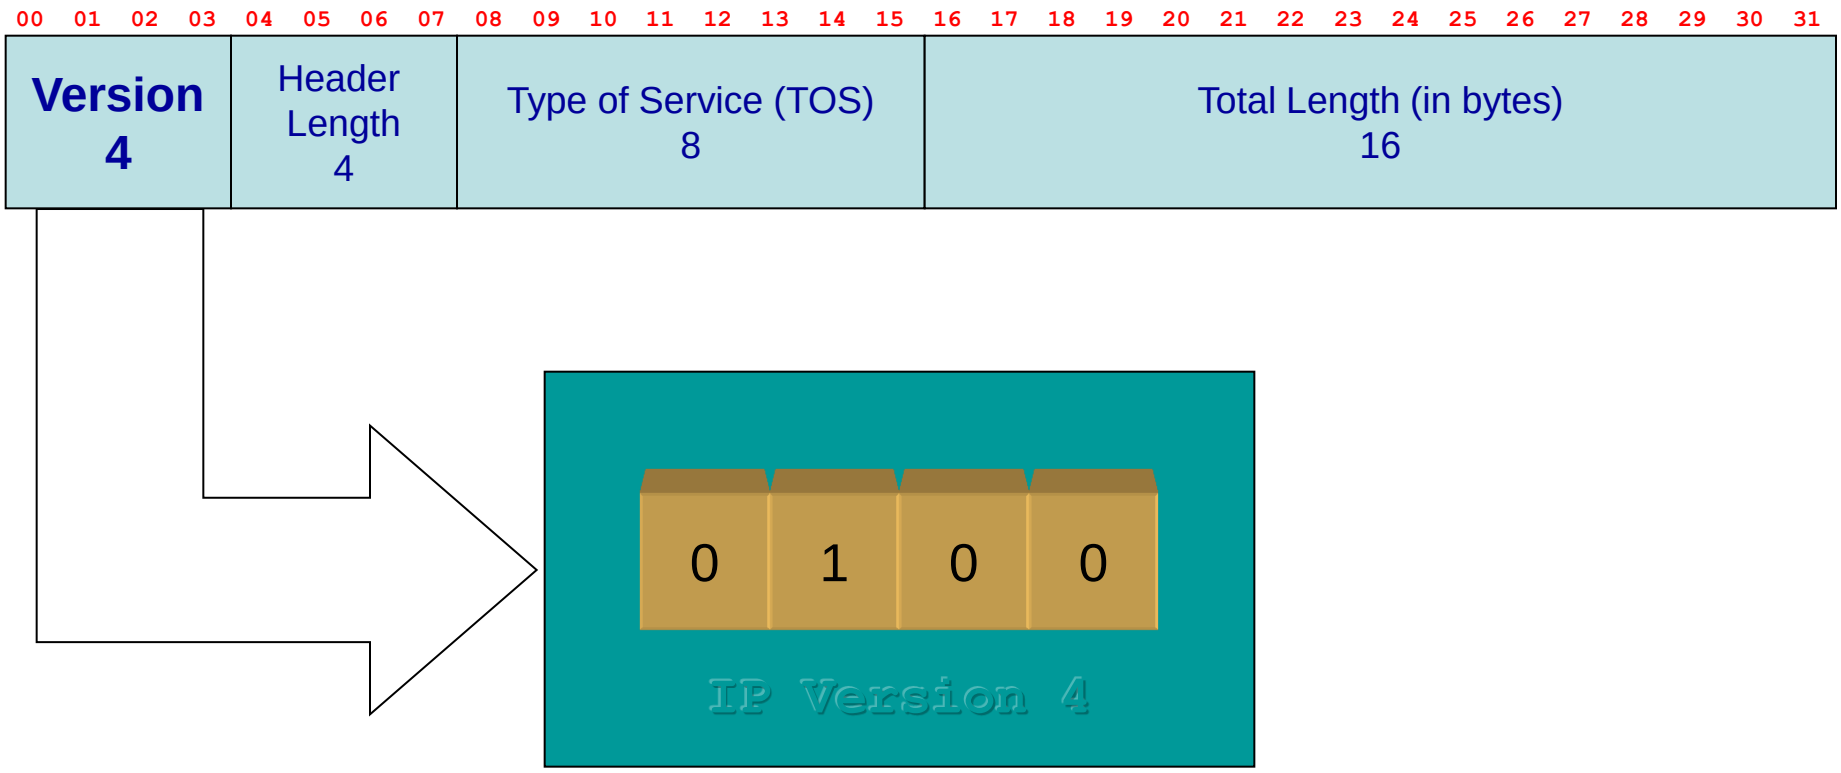

- Version field is 4 bits long
- It is the release version of IP

/network layer/ip/header/length

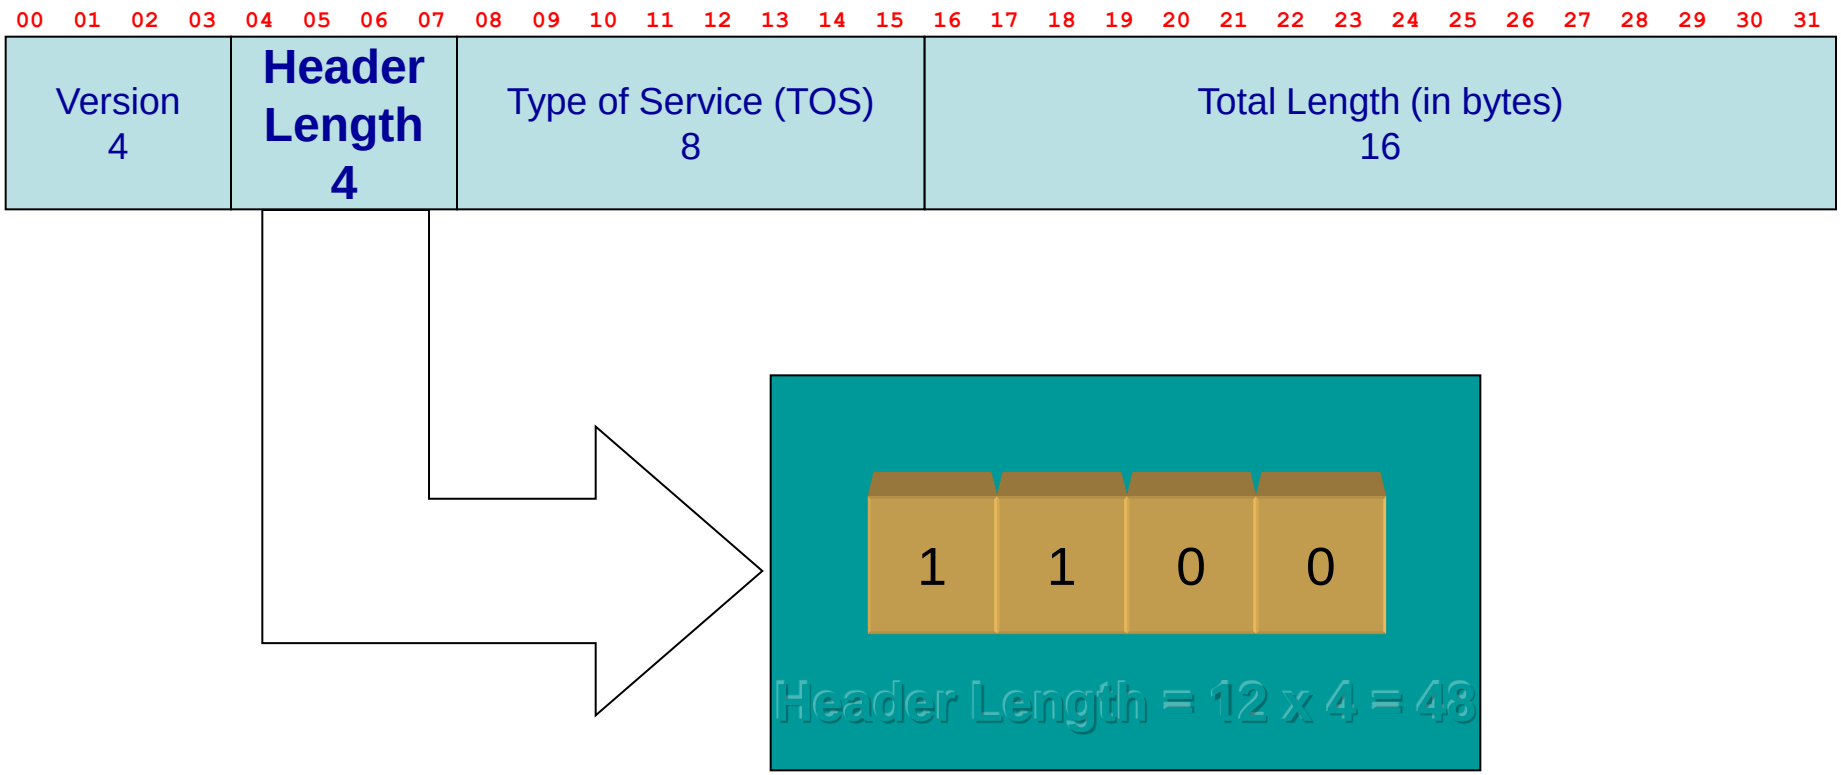

- Header length is number of 32 bit words
- It includes options
- Max = 60, Min = 20, it may need padding

/network layer/ip/header/TOS/precedence

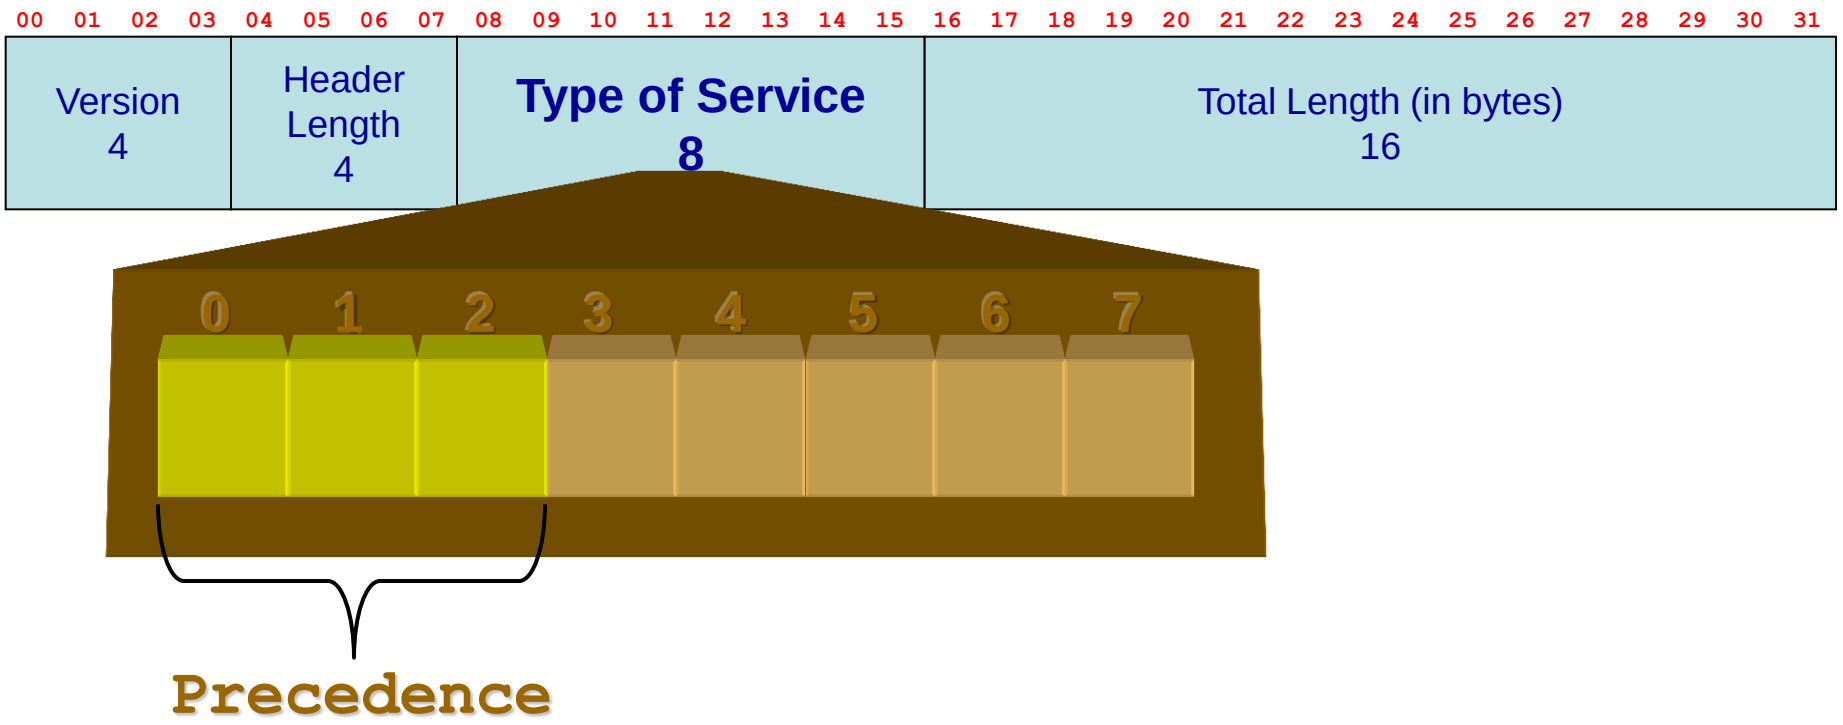

- A busy network can discard datagrams on the basis of its precedence
- 8 levels
- This field is used by router to handle congestion

/network layer/ip/header/TOS/bits

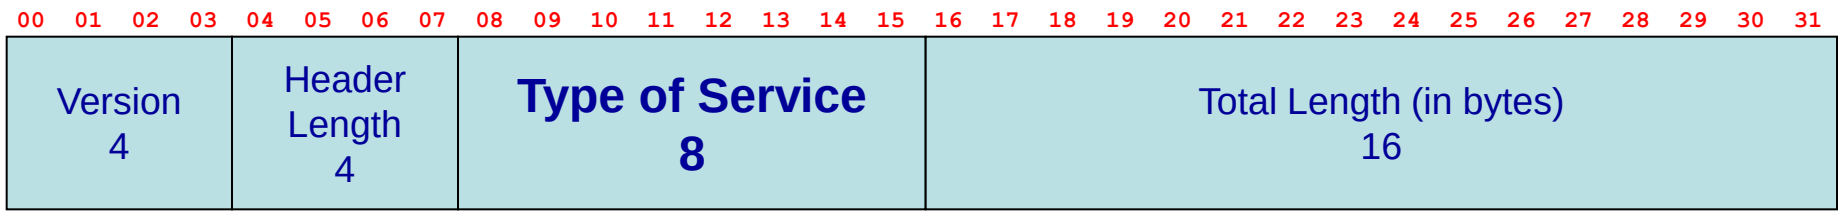

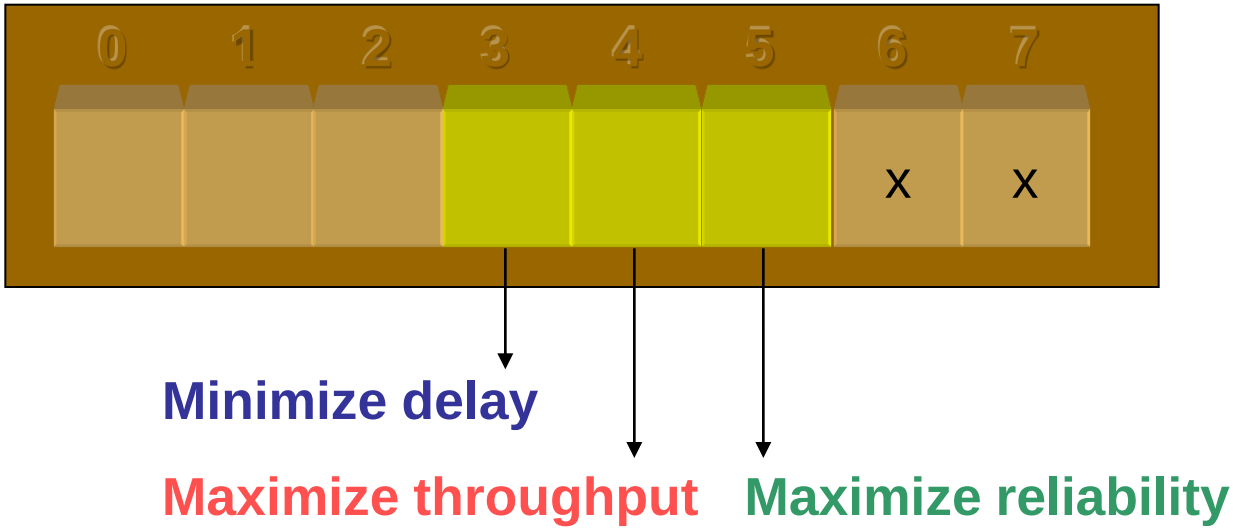

- Only 1 of these 3 bits can be turned on at a time
- All 3 bits set to 0 implies normal service

/network layer/ip/header/TOS/use

Application	Minimize Delay	Maximize throughput	Maximize reliability	Minimize monetary cost
Telnet / Rlogin	1			
FTP				
Control	1			
Data		1		
SMTP				
Command	1			
Data		1		
SNMP			1	
NNTP				1

/network layer/ip/header/TOS/ds

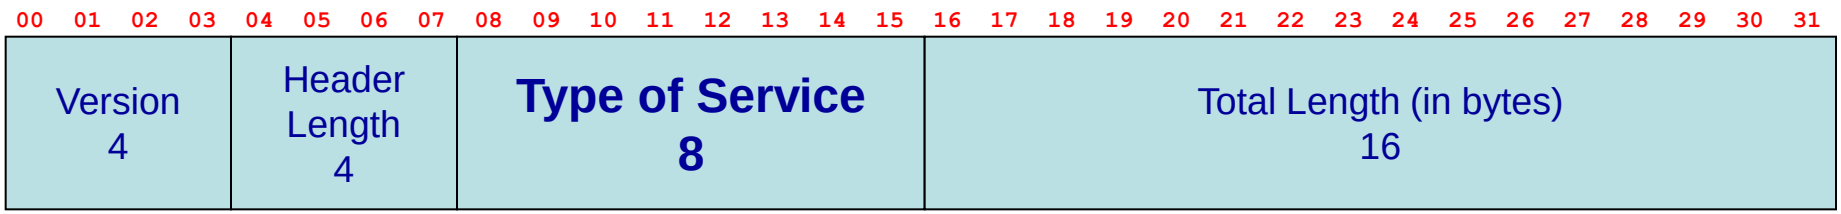

- RFC 2474 redefines the first six bits of the *TOS* field to support a technique called *Differentiated Services (DS)*

/network layer/ip/header/total length

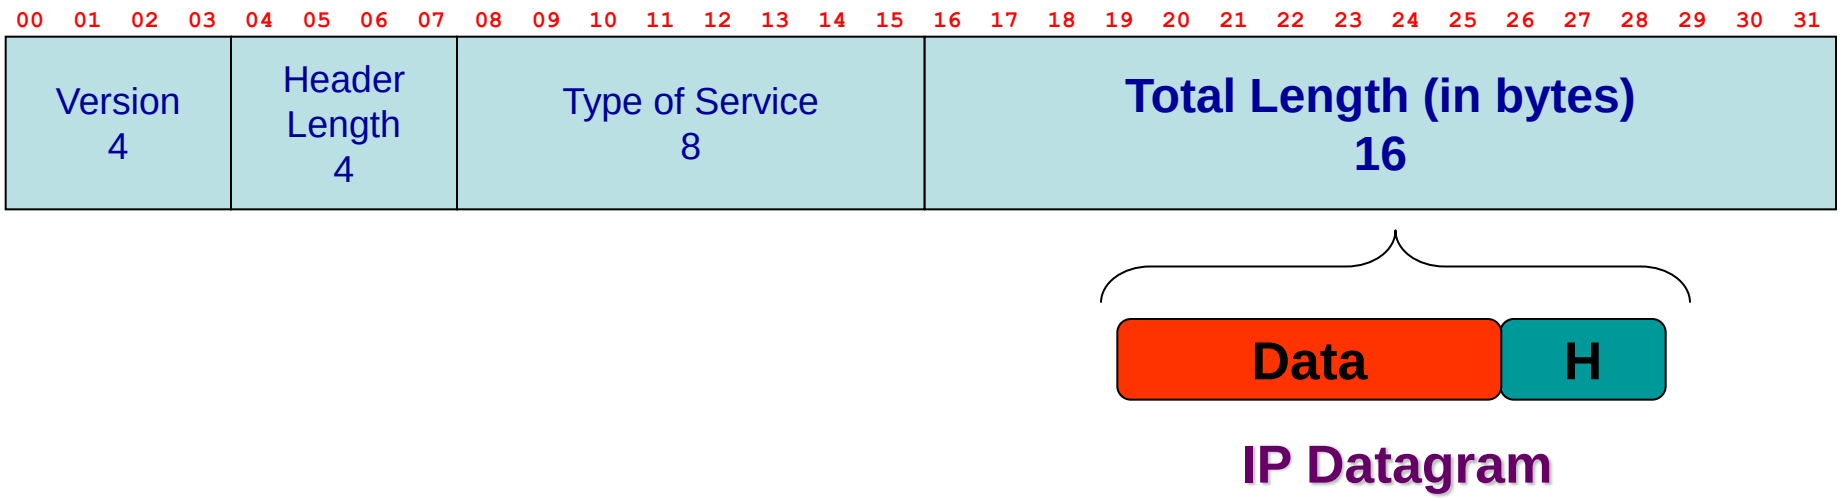

- Max size of an IP datagram = 64k
- Total length can change if a datagram is broken into multiple fragments

/network layer/ip/header/identification

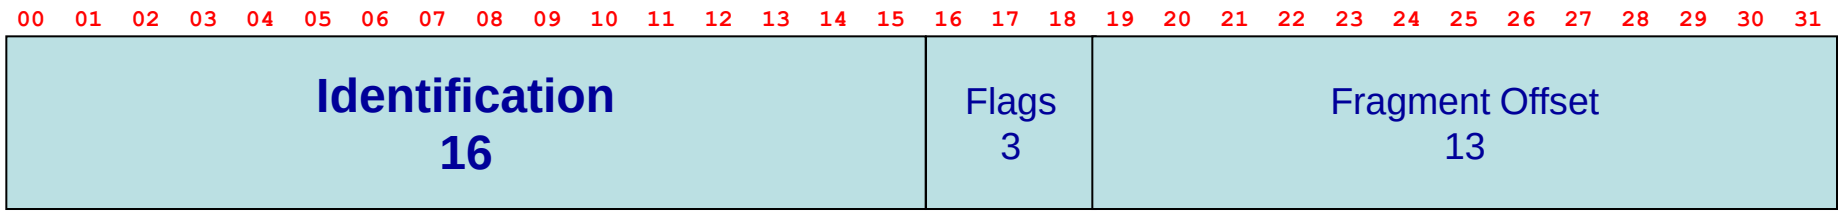

- Contains a unique value for each IP datagram, normally incremented by 1
- If a datagram is broken into multiple fragments, then this number is copied into each of those fragments
- A fragment is a datagram with its own IP header and is routed independently of any other datagrams

/network layer/ip/header/flags

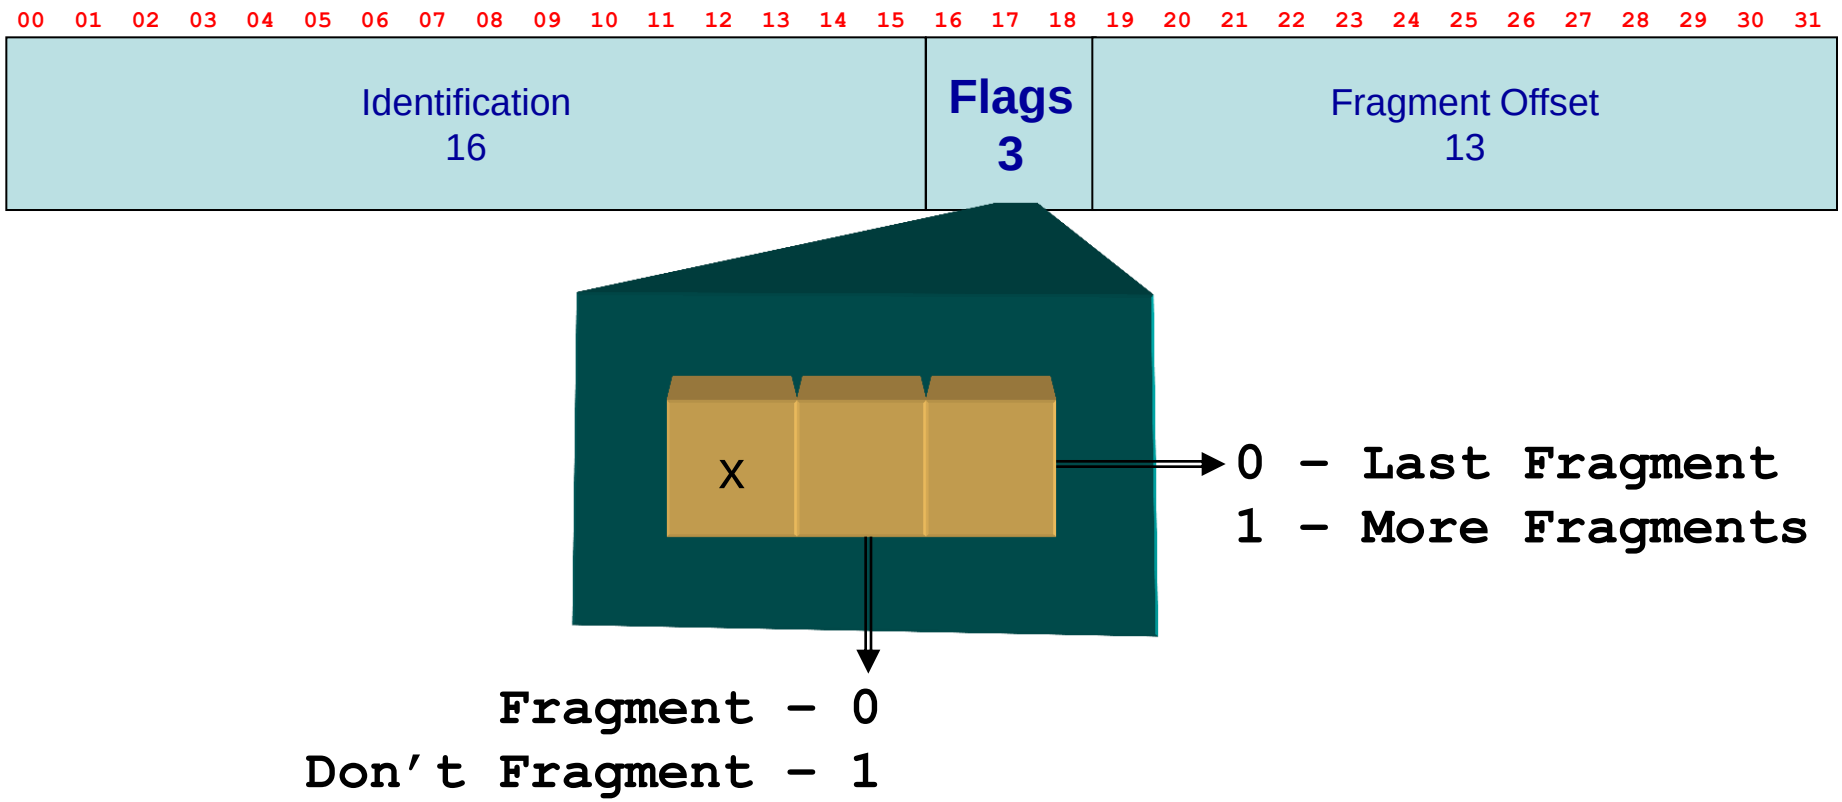

- Fragment offset contains offset of a fragment from the beginning of original datagram. It is specified in units of 8 bytes (64 bits)

/network layer/ip/header/fragmentation

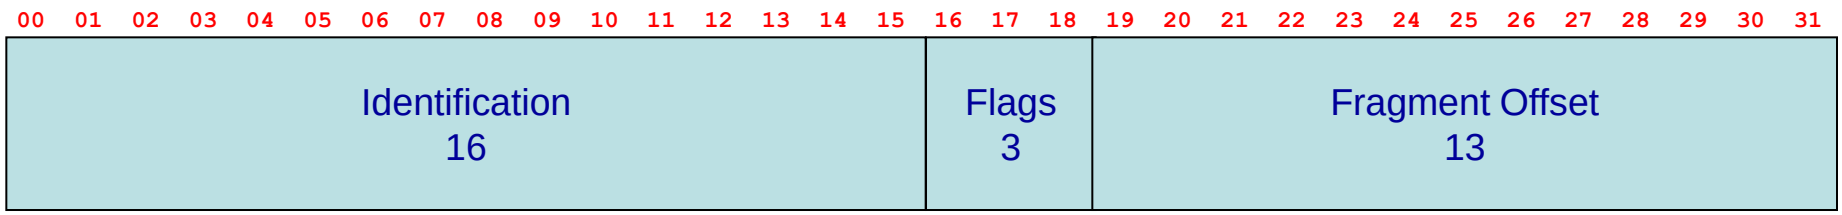

- Link layer imposes an upper limit on the size of the frame that can be transmitted
- IP queries and obtains link layer's MTU
- IP compares MTU with the datagram size and performs fragmentation if necessary
- Fragmentation may be done either by sending host or by an intermediate router
- Fragments are assembled only at the destination host

/network layer/ip/header/fragmentation/example

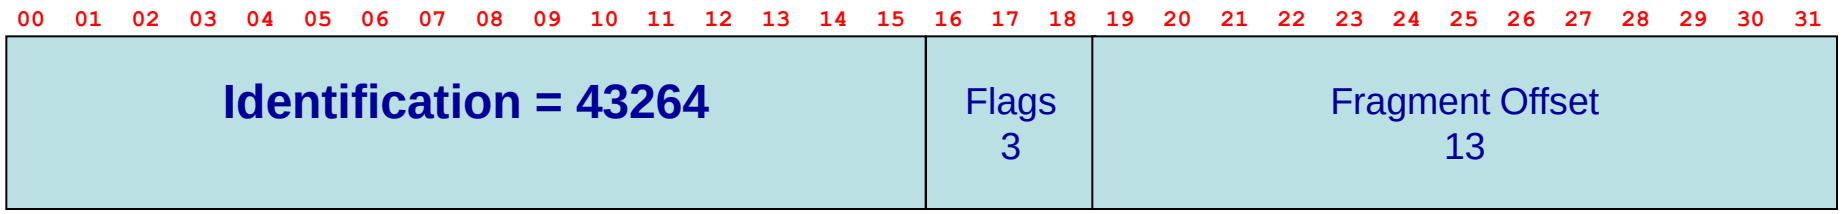

Total Length = 3020 bytes

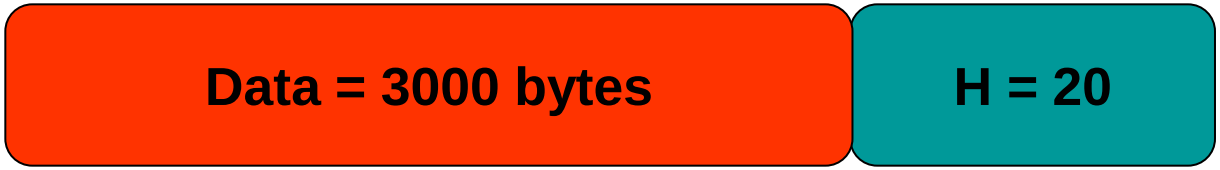

Data	Hdr
1480	20
1480	20
40	20
3000	**60**

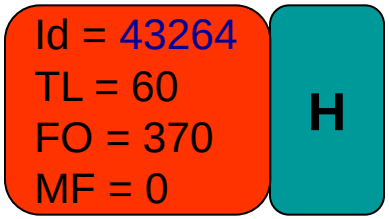

Actual FO: 2960 / 8

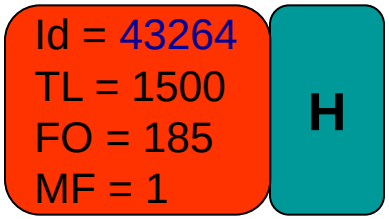

Actual FO: 1480 / 8

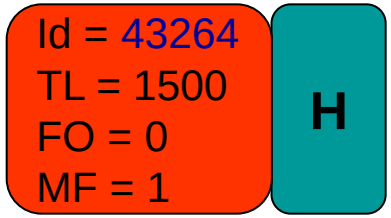

/network layer/ip/header/fragmentation

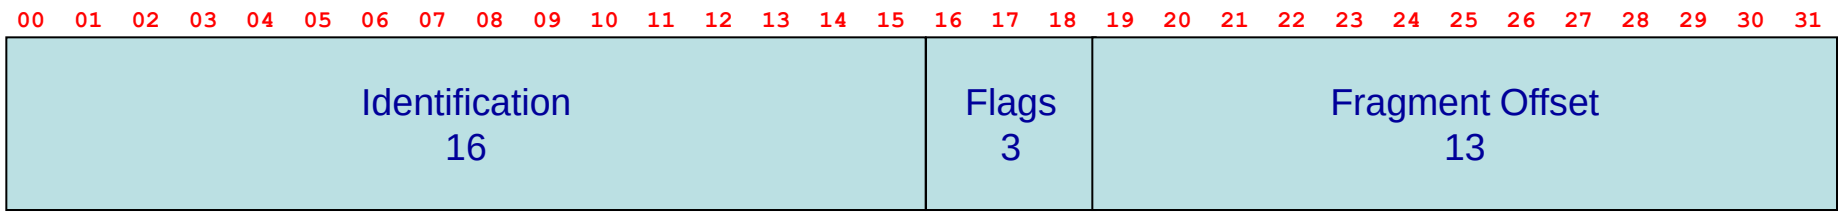

- Fragmentation and reassembly is transparent to transport layer
- If one fragment is lost, entire datagram has to be retransmitted
- If *'don't fragment'* bit is set, IP router will not fragment that datagram
- Fragmentation can cause performance degradation

/network layer/ip/header/TTL

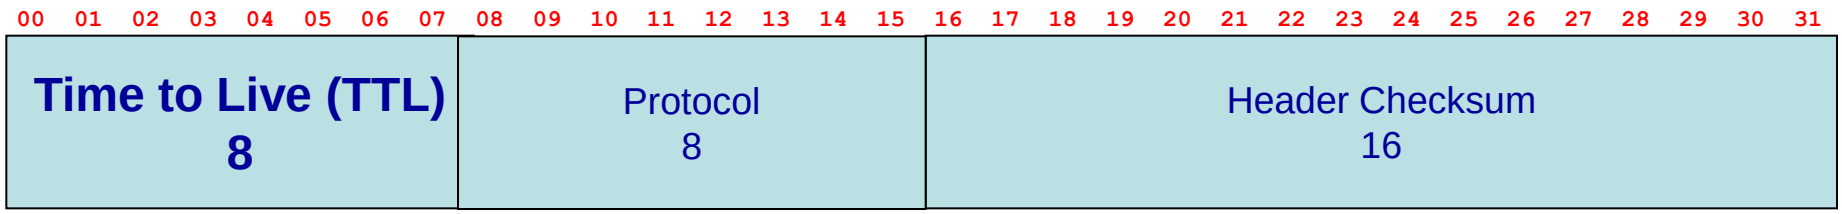

- IP does not know the complete route to any destination
- IP routing is done on a hop-by-hop basis
- TTL is an upper limit initialized by the sender, on the number of routers through which a datagram can pass
- TTL is decremented by 1 by every router that handles the datagram
- When TTL = 0, datagram is thrown away and the sender is notified by an ICMP "Time exceeded" message

/network layer/ip/header/protocol

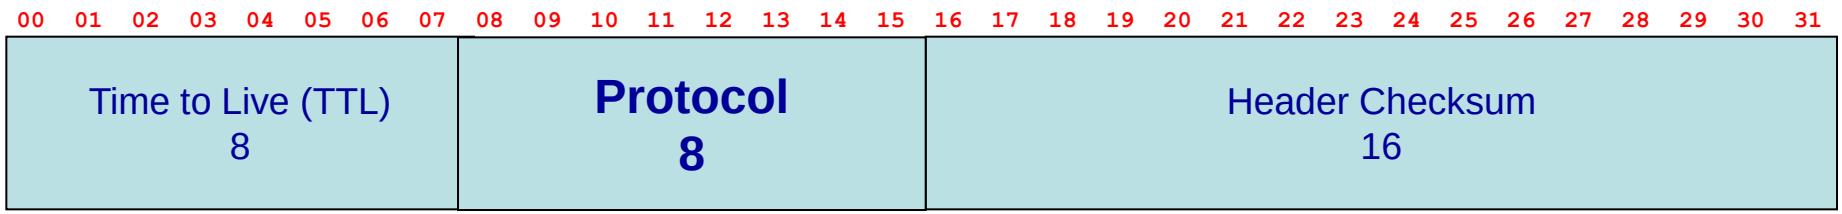

- Identifies upper layer protocol that gave the data to IP to send or is the intended recipient

ICMP = 1

TCP = 6

UDP = 17

/network layer/ip/header/checksum

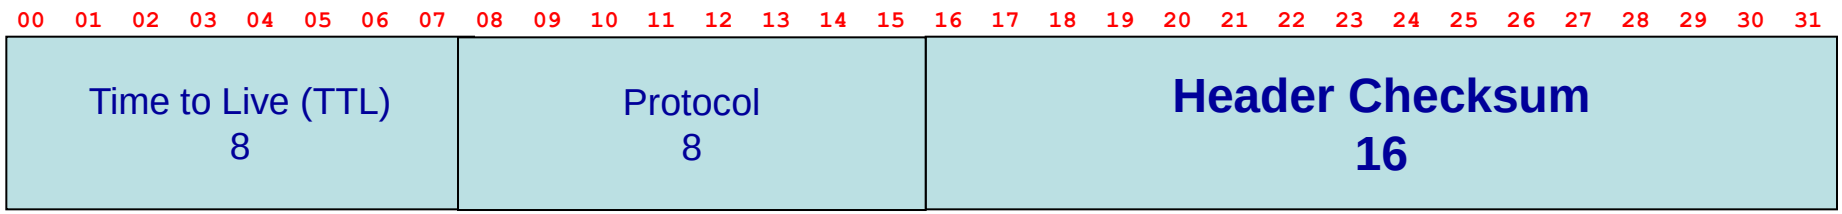

- Header checksum is calculated by 16-bit one's complement sum of the header
- The receiver of the datagram cross-checks integrity of the header by re-computing the checksum of the header and comparing it with the stored checksum
- If it does not match, IP discards the received datagram
- No error message is generated

/network layer/ip/header/address

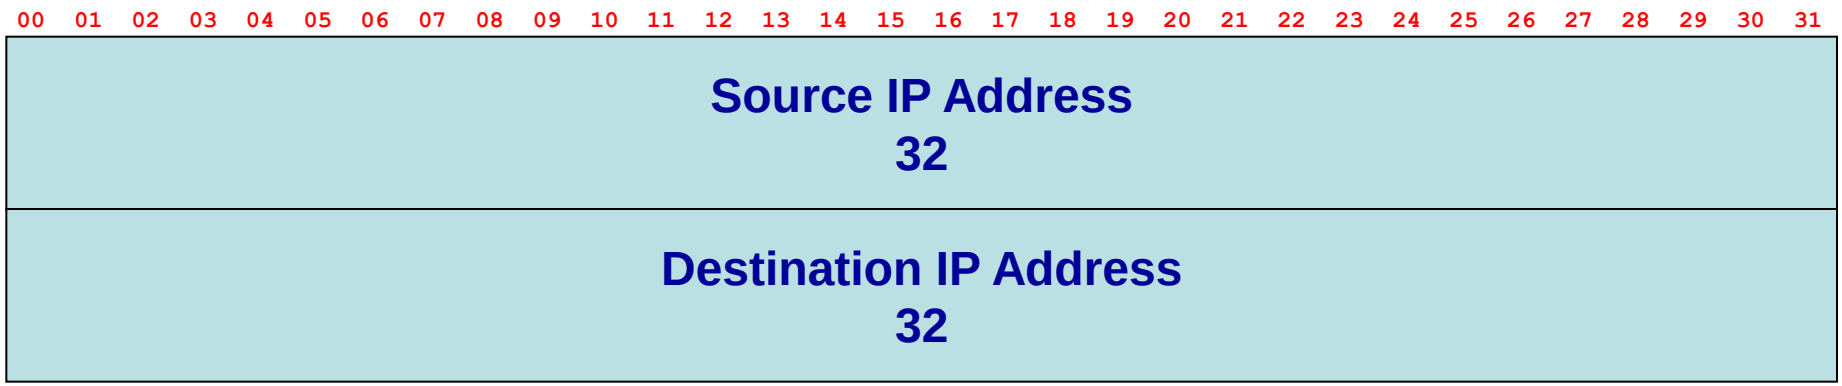

- 32 bit *valid* IP addresses

/network layer/ip/header/options

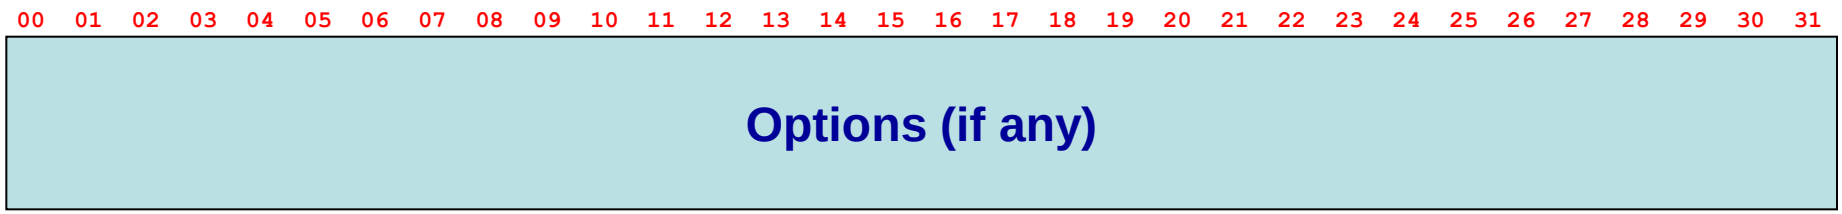

- It's a variable-length list of optional information for the datagram

- Options can't exceed 40 bytes

- Each option field has either 1 or 3 parts

  Type: 8 bits, identifies type of option

  Length: 8 bits, length of total option

  Data: variable length, applicable to option

/network layer/ip/header/options/list

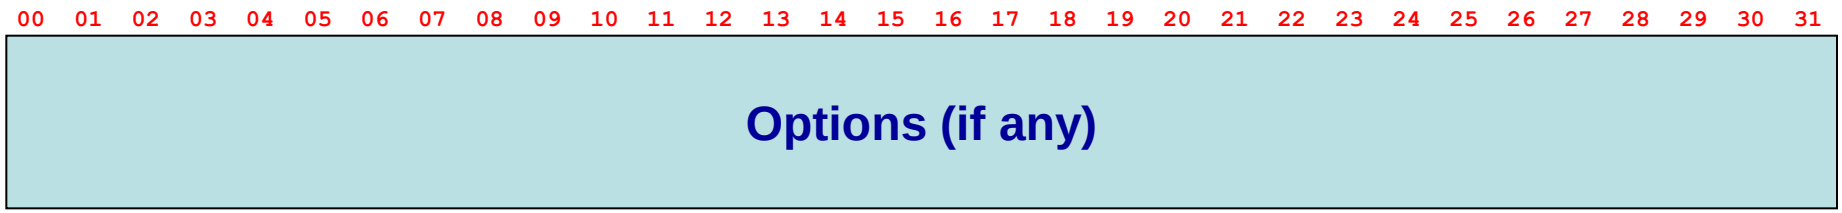

- Security and handling restrictions (for US military)
- Record route (have each router record its IP address)
- Timestamp (have each router record its IP address and time)
- Loose source routing (specifying a list of IP addrs that must be traversed by datagram)
- Strict source routing (only the addrs in the list can be traversed)

/network layer/ip/header/options/record route

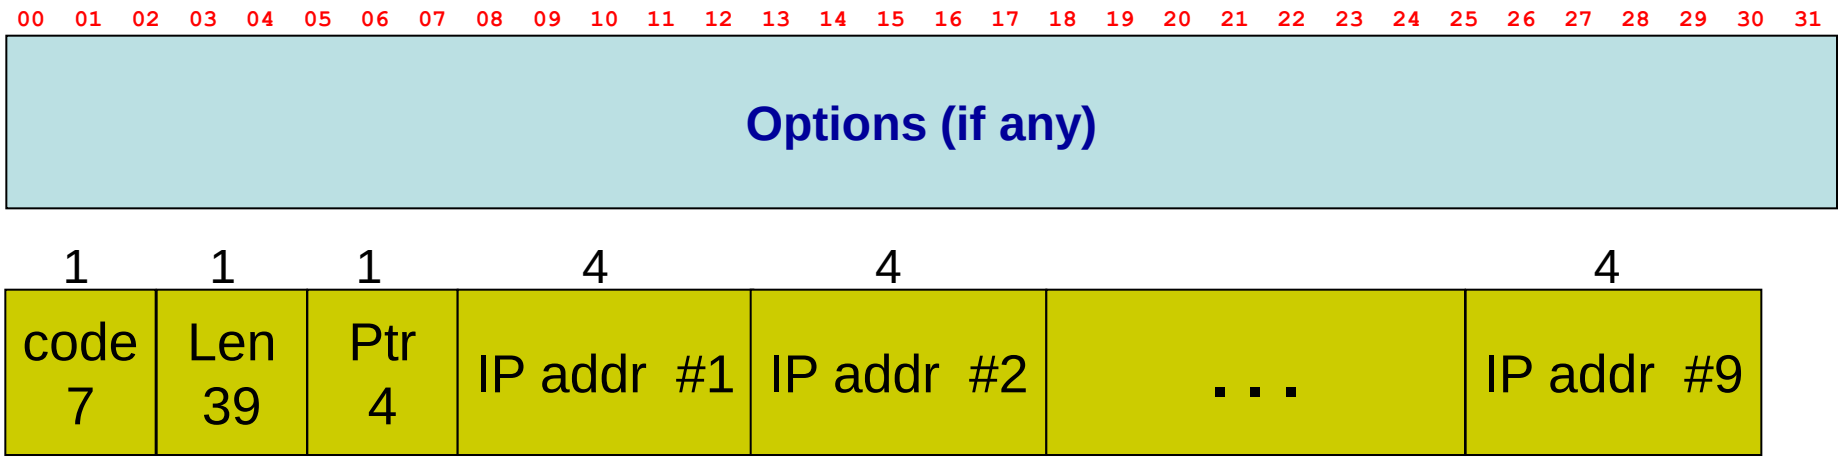

- Every router that handles the datagram with above option set, adds its own IP addr to a list in *options* field
- This feature is used for knowing the path (addrs of all the routers) through which the datagram passed on its way to the destination host
- Used by *ping* utility when used with –r parameter

/network layer/ip/header/options/timestamp

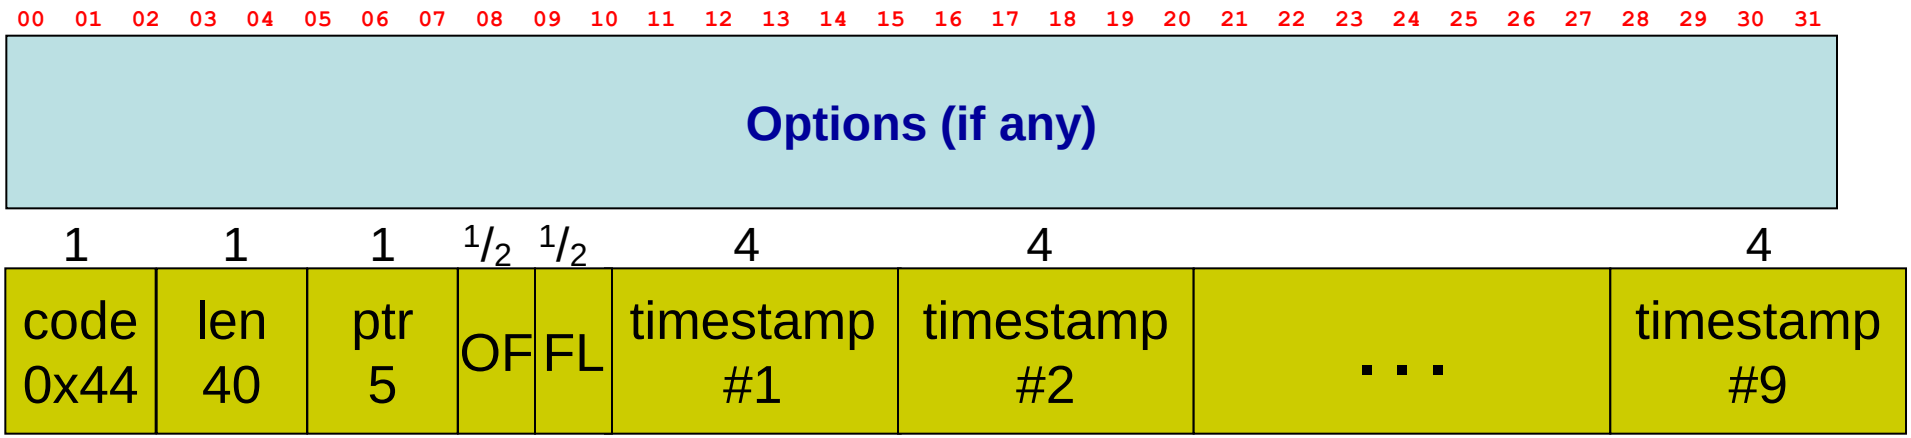

- Timestamp is the number of milliseconds past midnight of a system (can also be some other format)
- If a router can't add timestamp due to shortage of space, it increments *overflow* field by 1

flags	description
0	Record only timestamp
1	Each router records IP Address + Timestamp
3	A router records its timestamp only if its IP addr is in the list initialized by the sender

/network layer/ip/header/options/source routing

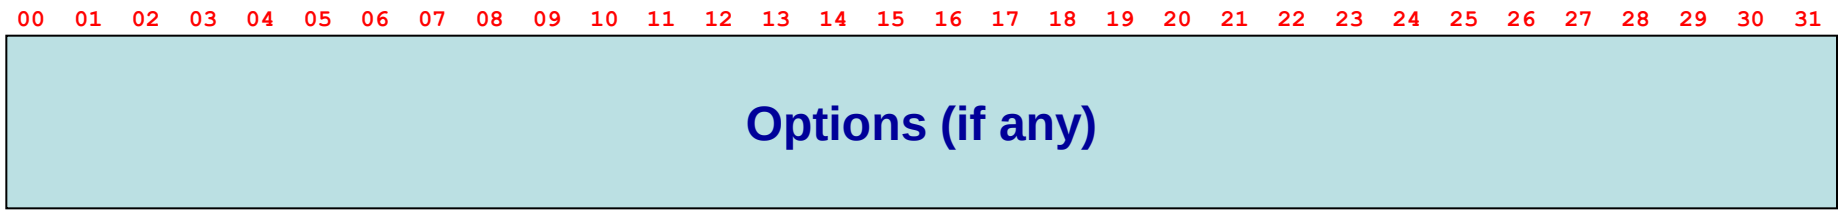

- *Strict:* The sender specifies exact path that the IP datagram must follow. Code = 0x83

- *Loose:* As above, except that the datagram can also pass through other routers between any two addresses in the list. Code = 0x89

/network layer/ip/header

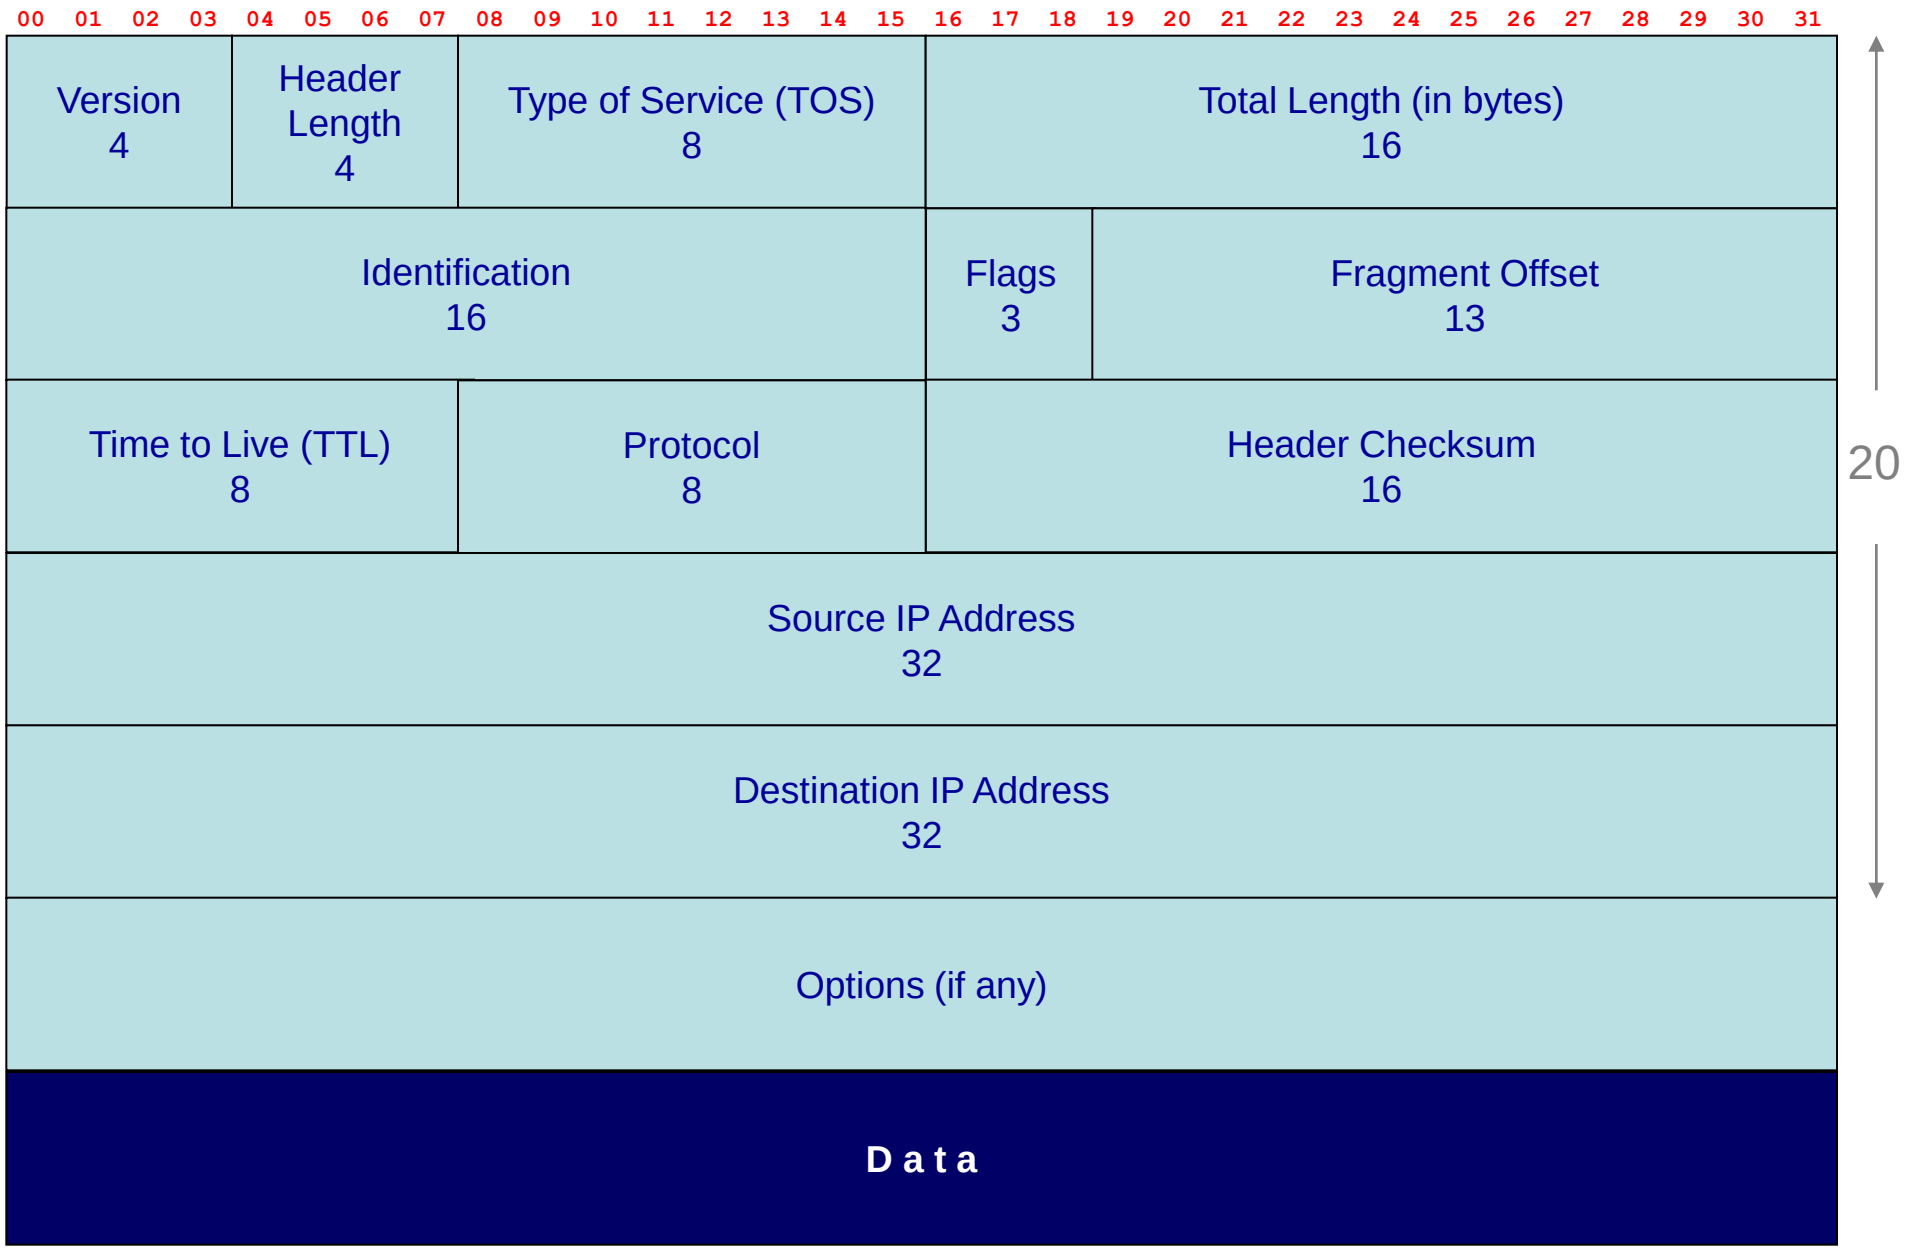

/network layer/ip/assignment

- ifconfig
  - Study and understand its output for various interfaces on the host
- netstat
  - Print MTU of each interface on your host
  - Print routing table
- tcpdump
  - Trap and display IP packets
  - Display only header part of IP packet. Study this header
- traceroute
  - Study program operation
  - Find path between any two hosts using traceroute
  - Explain the output

/network layer/icmp

- Communicates error messages and other conditions that require attention
- These messages are used either by IP or by TCP/UDP
- All ICMP messages are encapsulated in an IP datagram

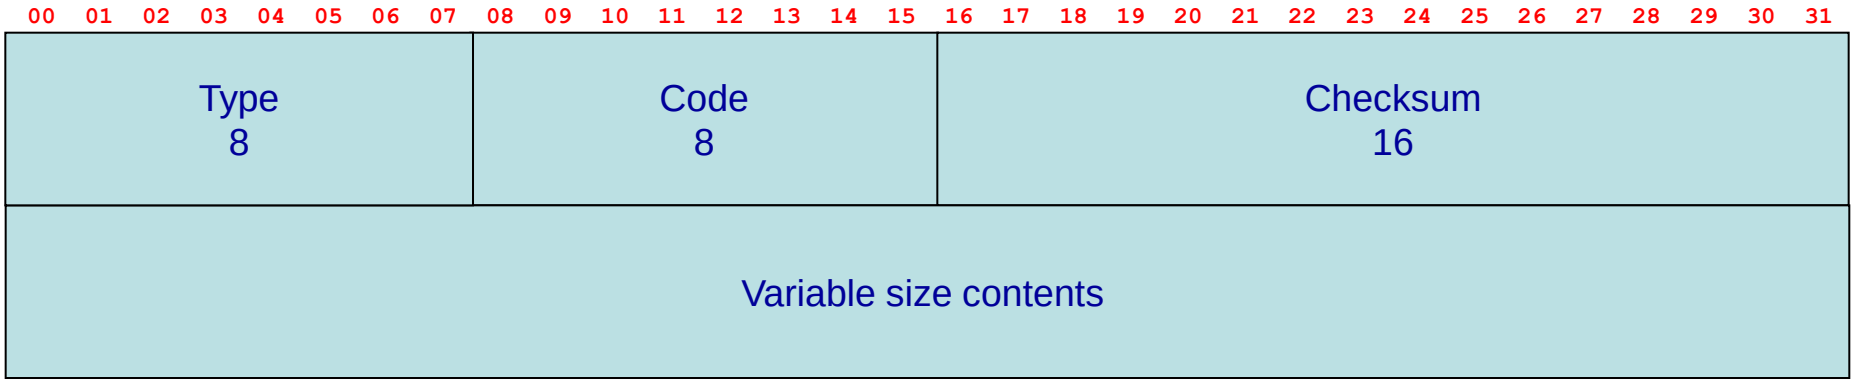

- ICMP error message always contains the IP header and first 8 bytes of the IP datagram that caused this error to be generated

/network layer/icmp/message types

Type	Code	Description	Query	Error
0	0	echo reply	•	
3	0	network unreachable		•
	1	host unreachable		•
	2	protocol unreachable		•
	3	port unreachable		•
	4	fragmentation needed…		•
4	0	source quench		•
8	0	echo request	•	
11	0	TTL = 0 during transit		•

/network layer/icmp/message types

Type	Code	Description	Handled by/msg
0	0	echo reply	user process
3	0	network unreachable	“no route to host”
	1	host unreachable	“no route to host”
	2	protocol unreachable	“connection refused”
	3	port unreachable	“connection refused”
	4	fragmentation needed…	“message too long”
4	0	source quench	kernel for TCP
8	0	echo request	kernel generates reply
11	0	TTL = 0 during transit	“Time exceeded”

/network layer/icmp/rules

- An ICMP error message is never generated in response to:
  - An ICMP error message
  - A datagram destined to an IP broadcast address
  - A datagram sent as a link-layer broadcast
  - A fragment (other than first) of a datagram
  - A datagram whose source address does not specify a single host

1. Use *ping* to learn more about ICMP

2. Find round-trip-time from a host to another host

3. Trap only *ping* echo and reply packets using *tcpdump*

4. Use ping to record route of a datagram from host to destination host
5. Use ping to record timestamps of all routers that a datagram passes through while reaching its destination host

Questions ?